

BRUNO SCHEUFLER  
[AT]GOPHERCON  
UK → 2025



INNGEST  
EOCOM  
EST...2021  
©2025

INNGEST

## ■ Welcome

---

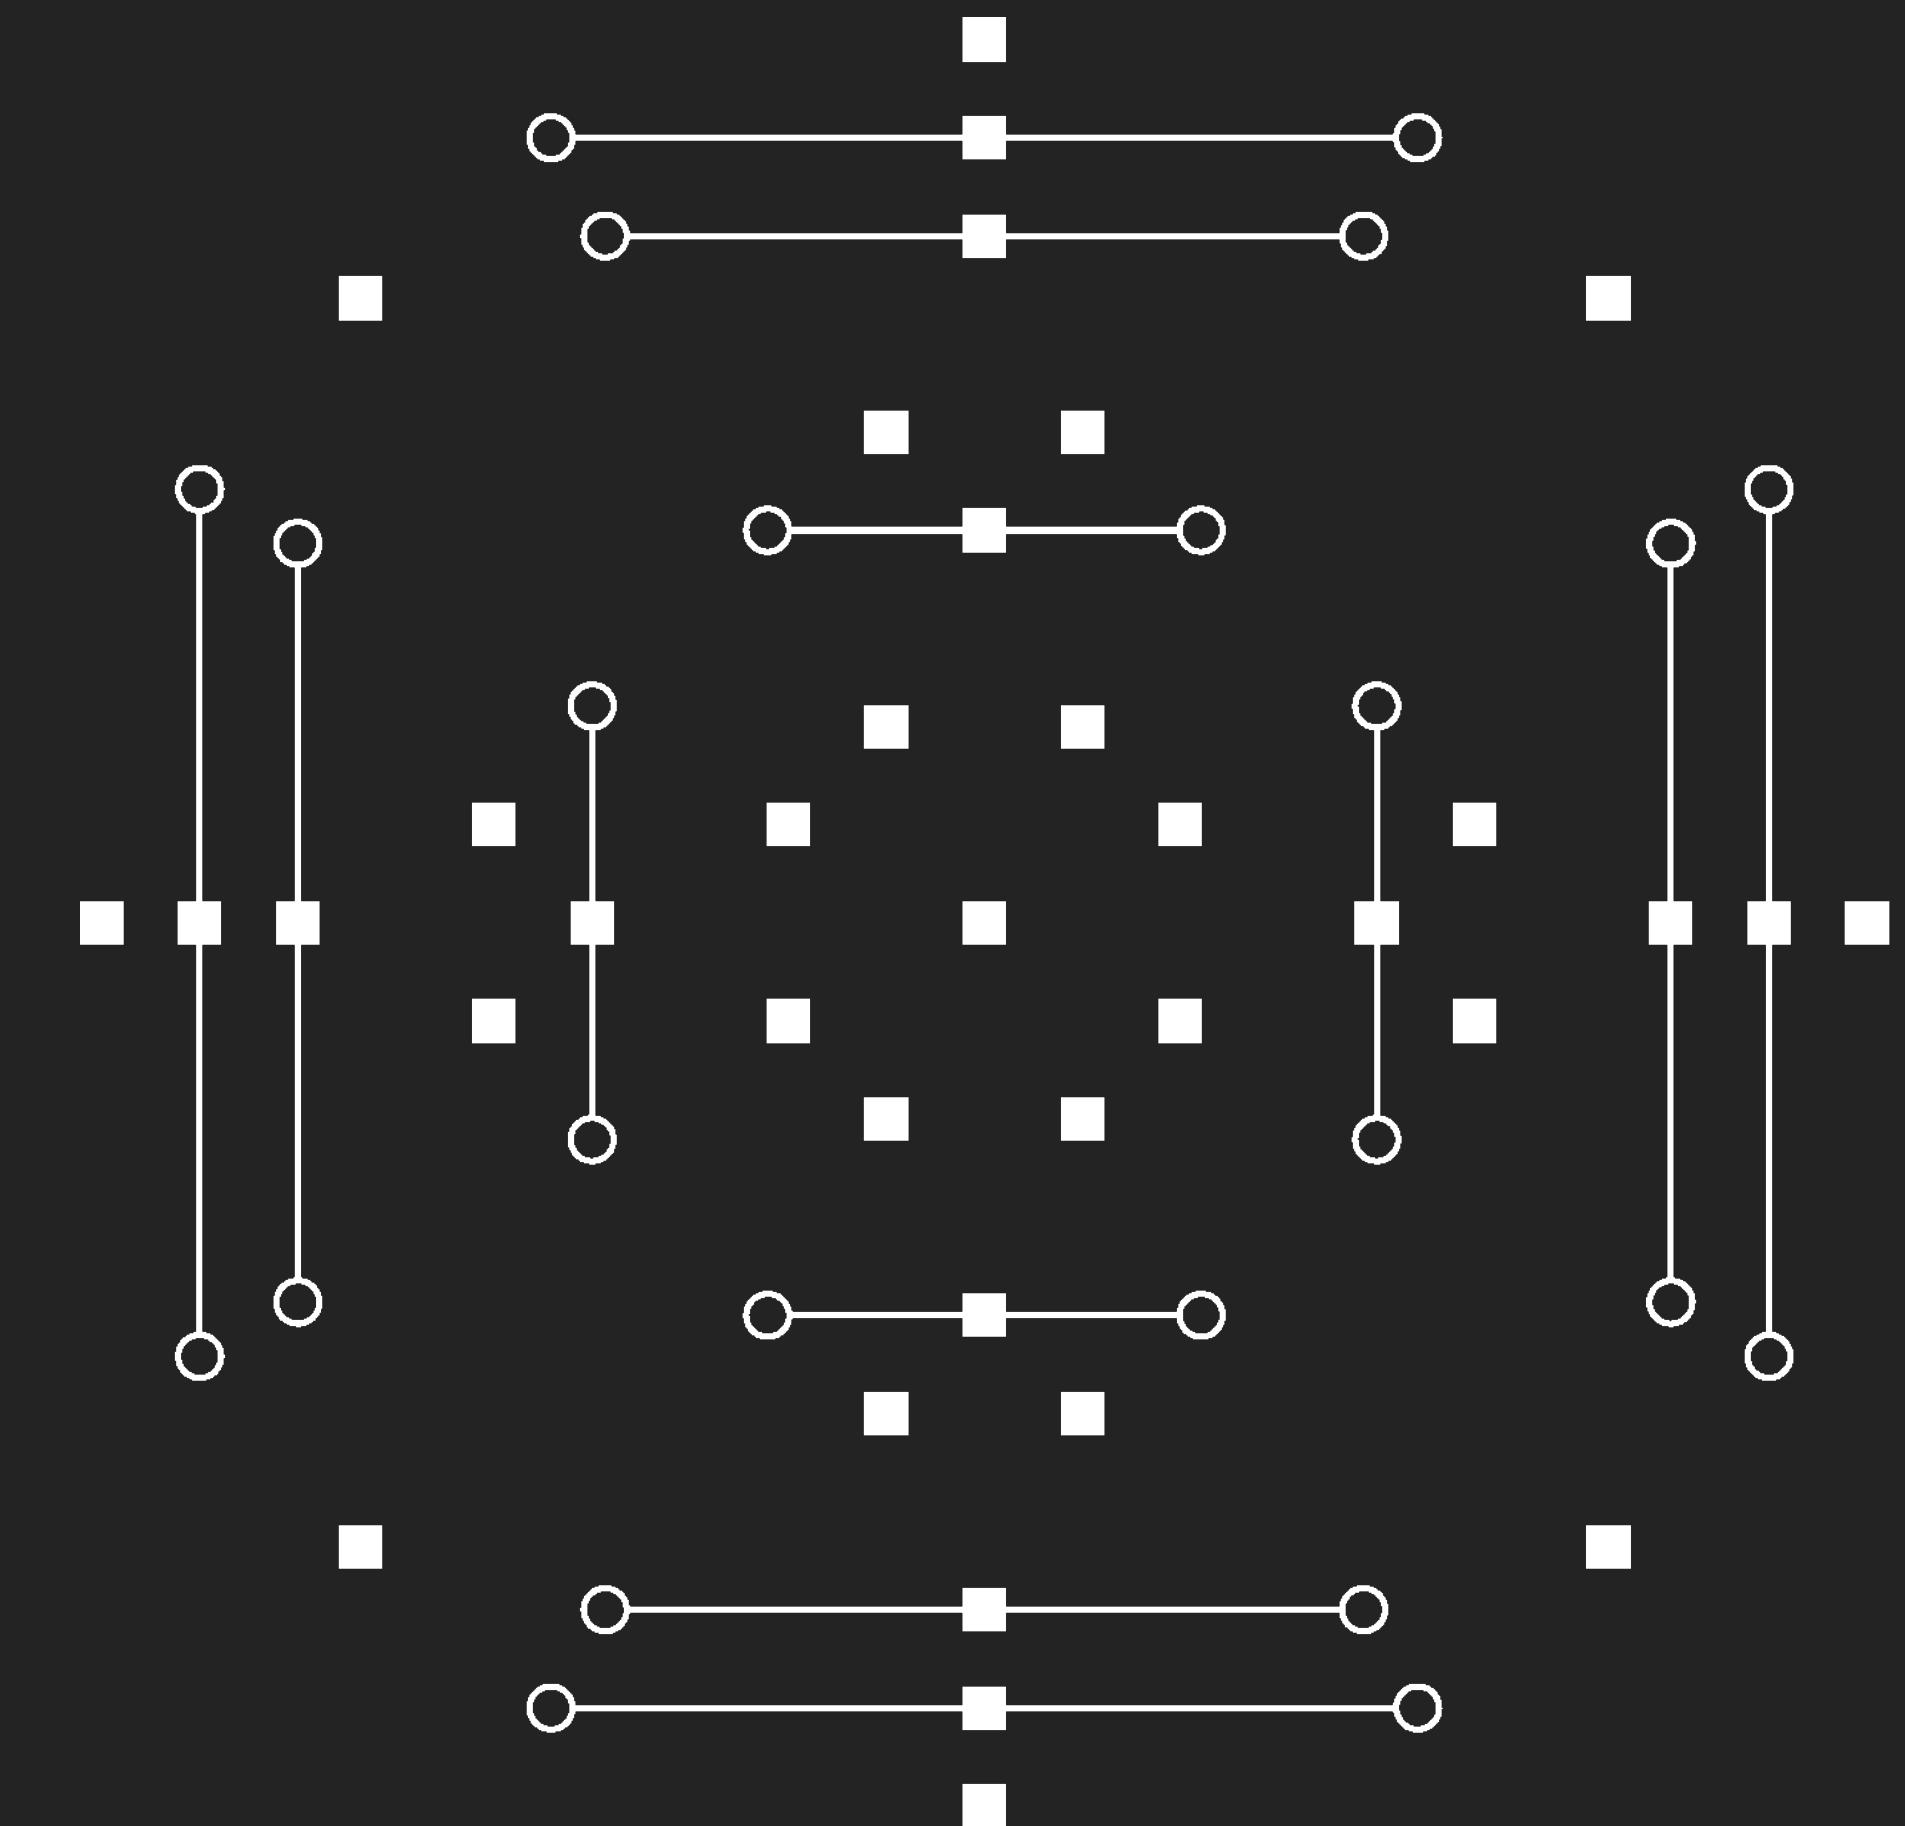
- ↳ Please move up to the front [↑]
  - ↳ We'll begin shortly...
- 

BRUNO SCHEUFLER  
[AT]GOPHERCON  
UK → 2025



INNGEST  
ECOM  
EST...2021  
©2025

# ■ Building a framework for reliable data migrations in Go



- Show of hands

- .

- .

- [01]

└[::/]└[ :: ]┘[ ↴ :: ]┘

# Who has written a migration before?

- Show of hands

- .

- .

- .

- [02]

└ [:: ⊞] └ [ :: ] └ [ ⊞ :: ] └

# Who has run into issues rolling out a migration?

■ Show of hands

.

.

.

[03]

└[::/]└[ :: ]└[ ↴ :: ]

Who has broken production  
with a migration?

# A story before we get started



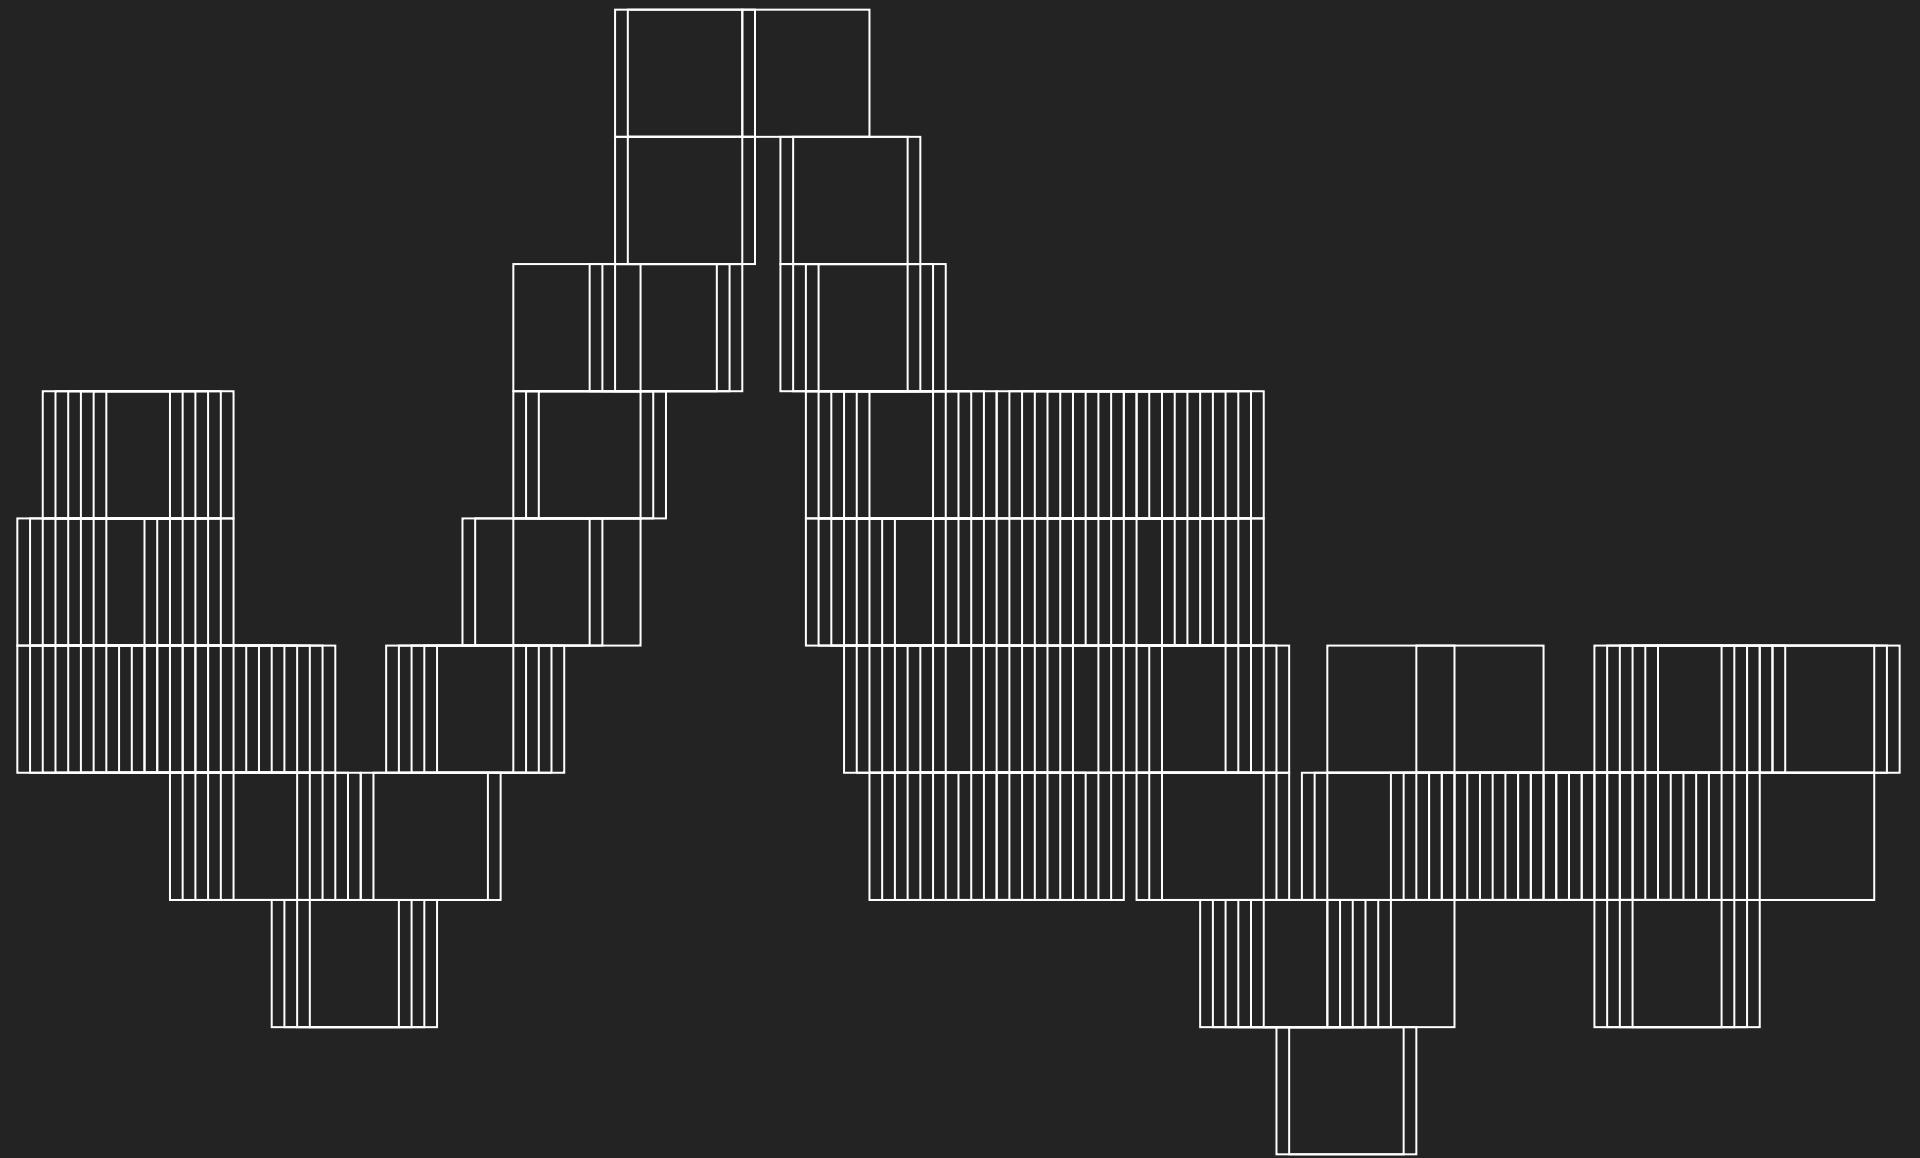
# Hello :) I'm Bruno

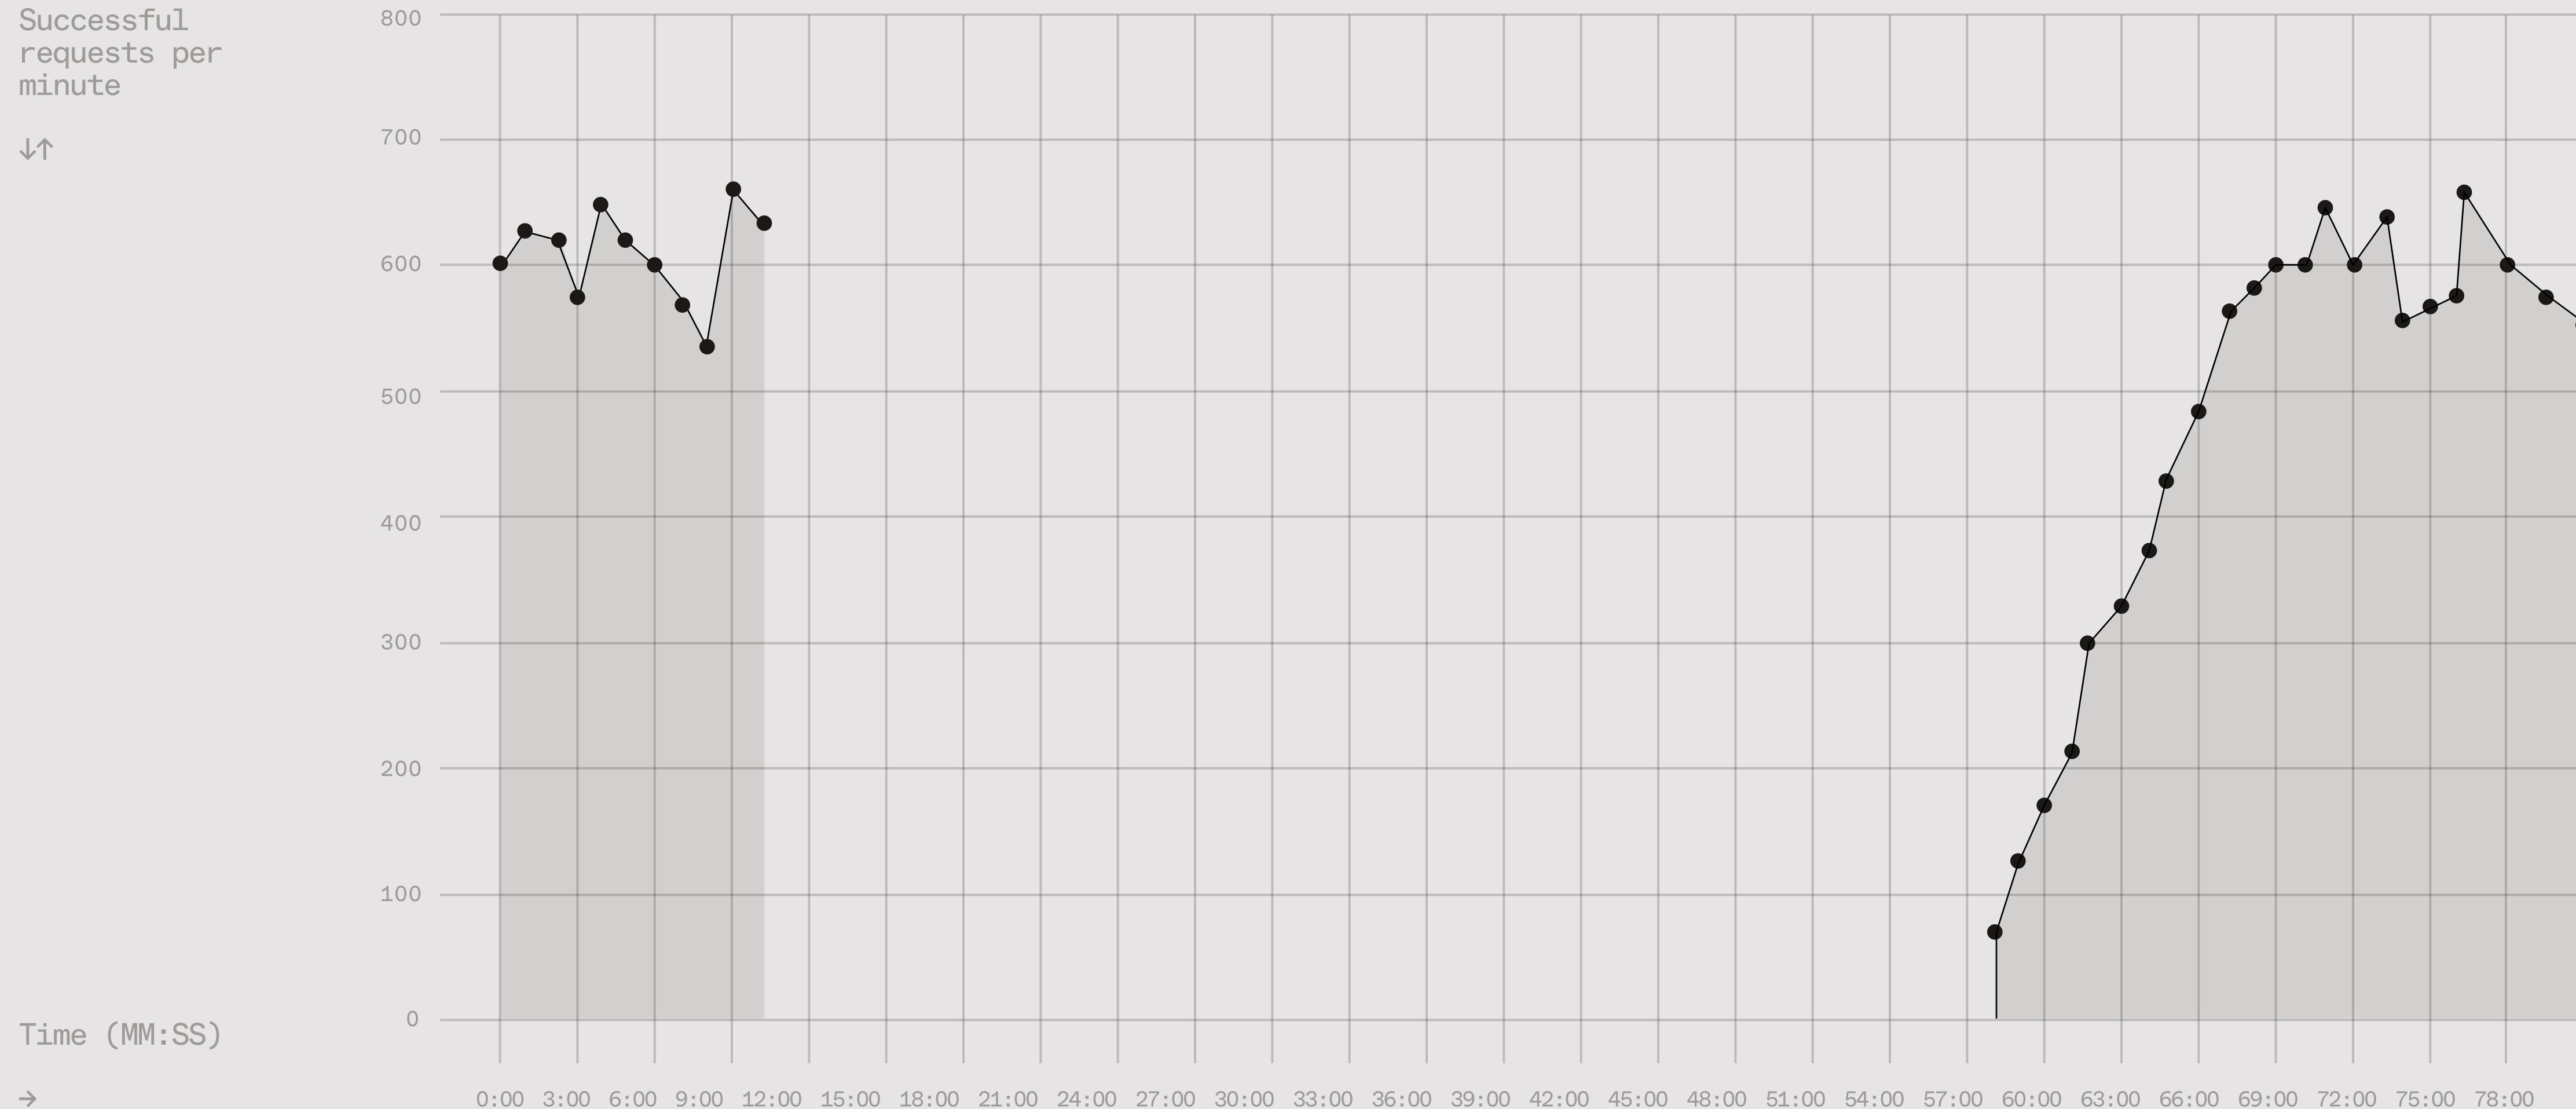
- ↳ Software Engineer
- ↳ @Inngest



■ 14 August/25

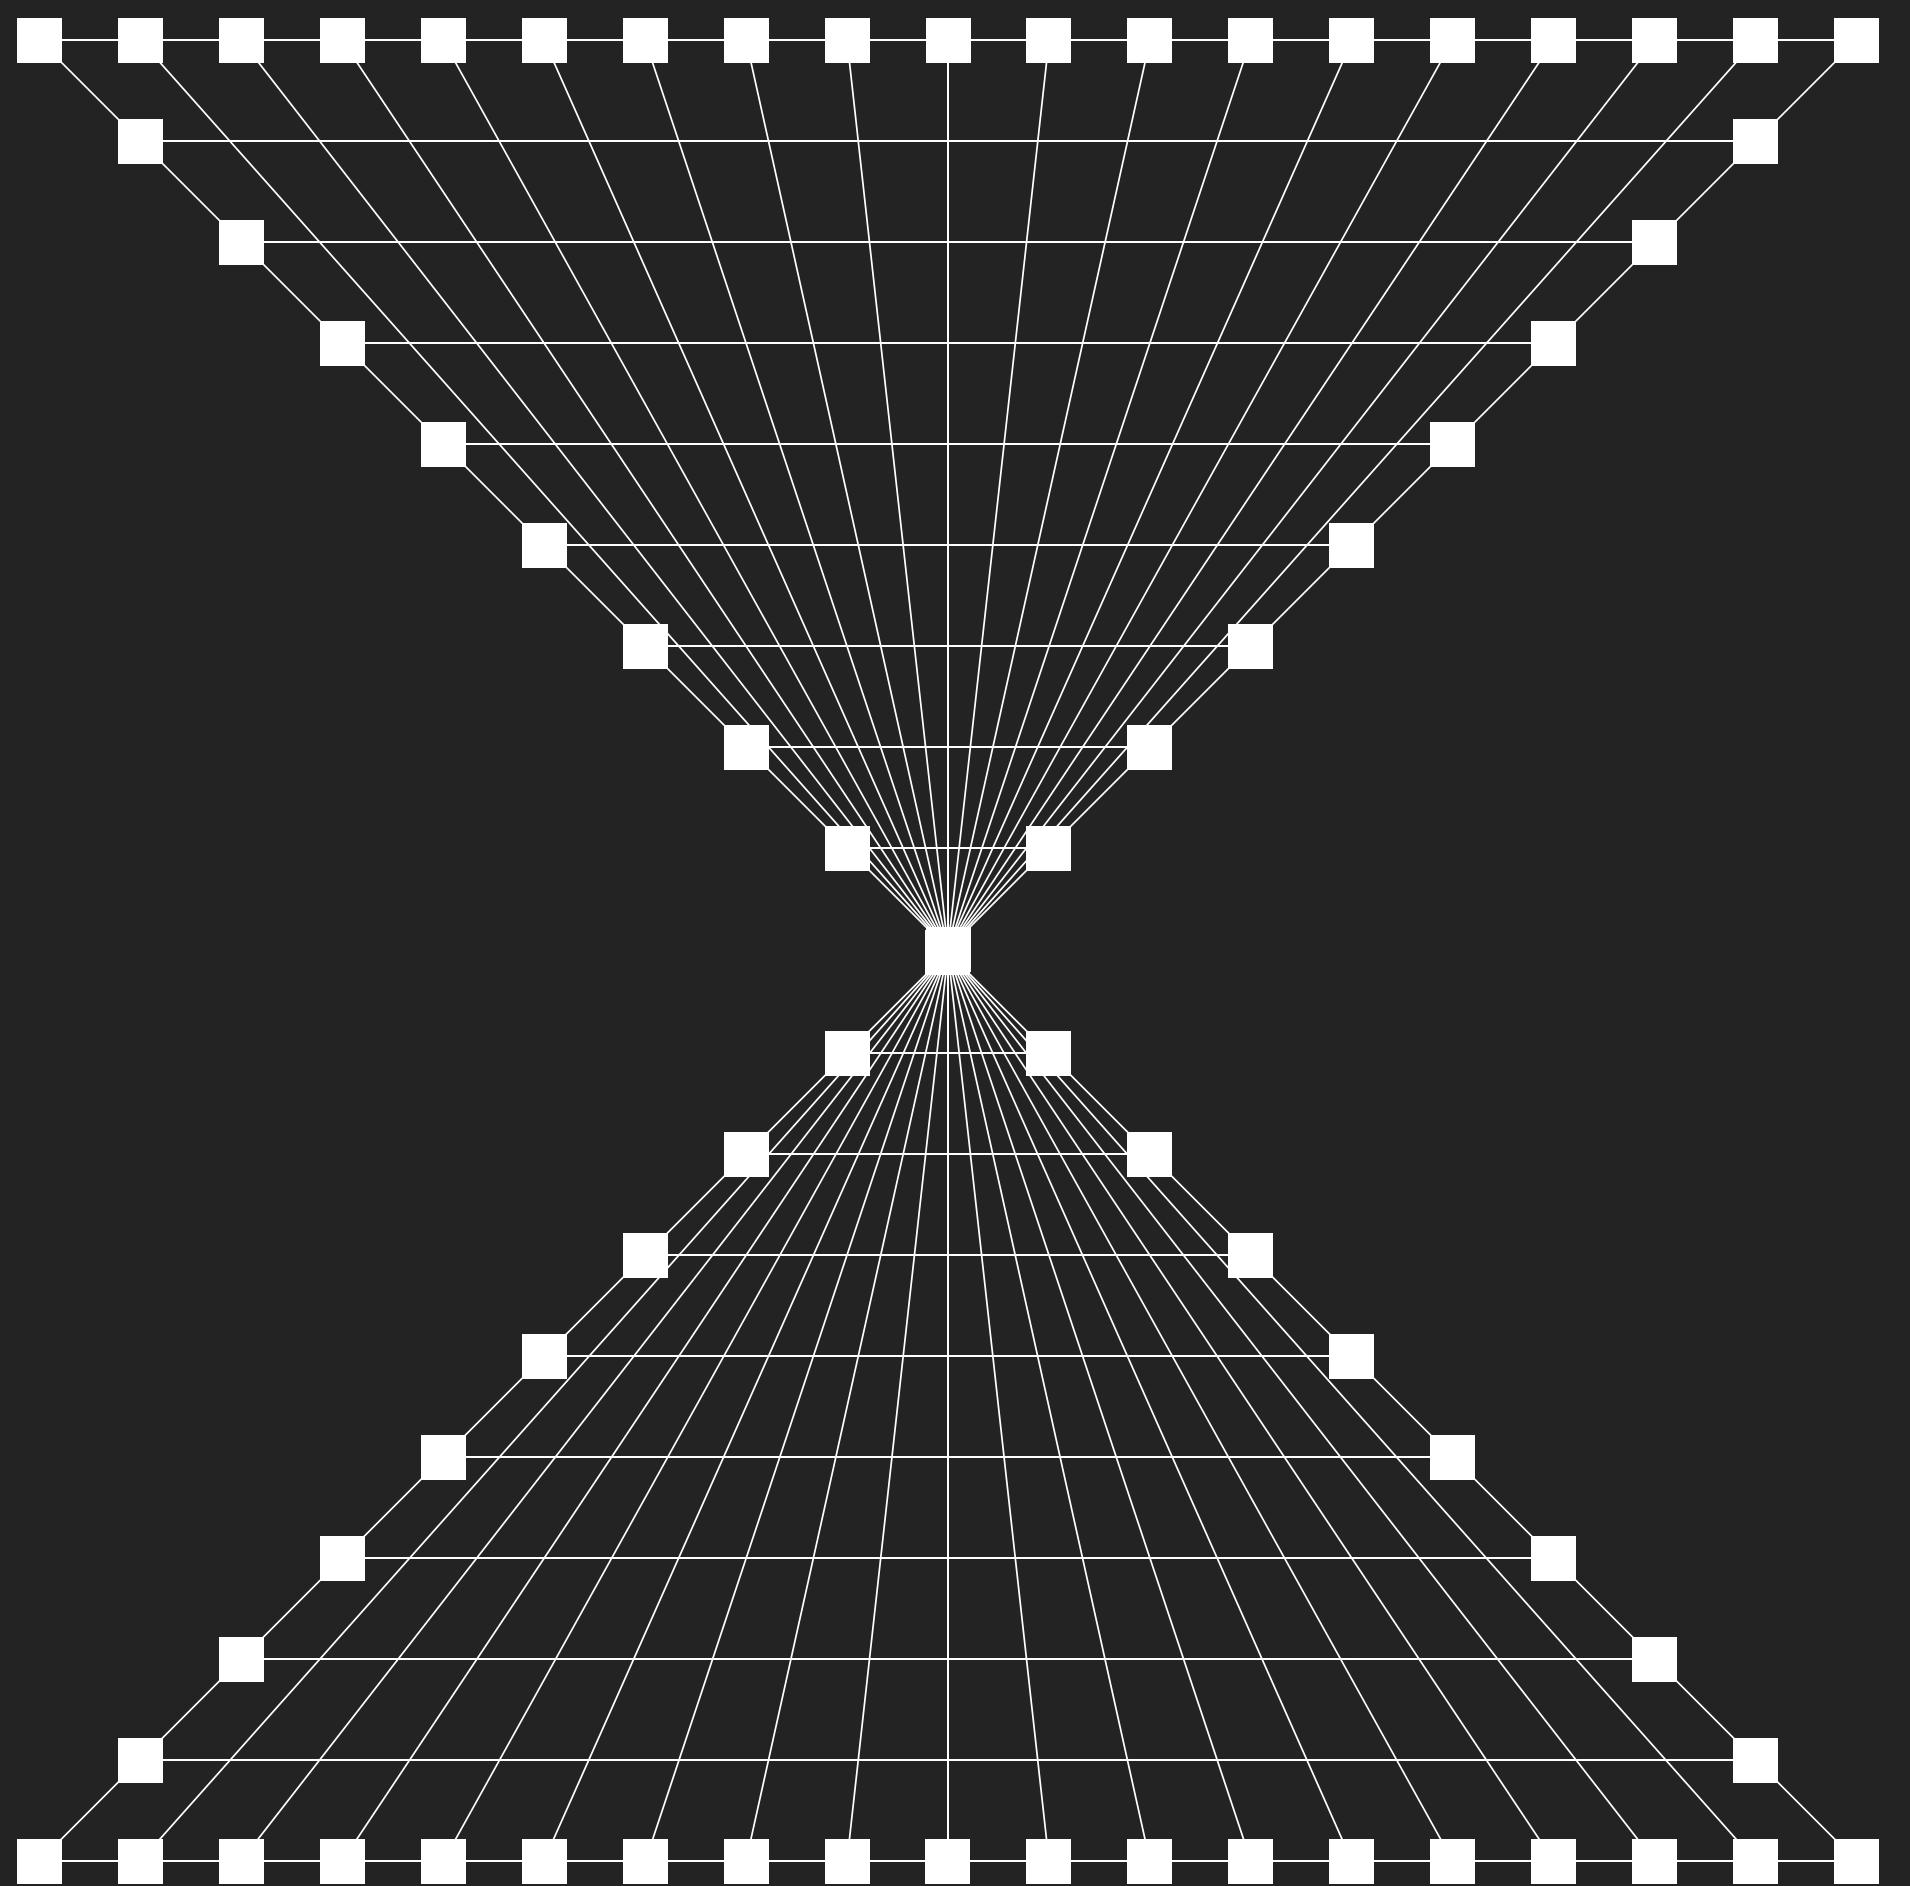
# How not to run a migration



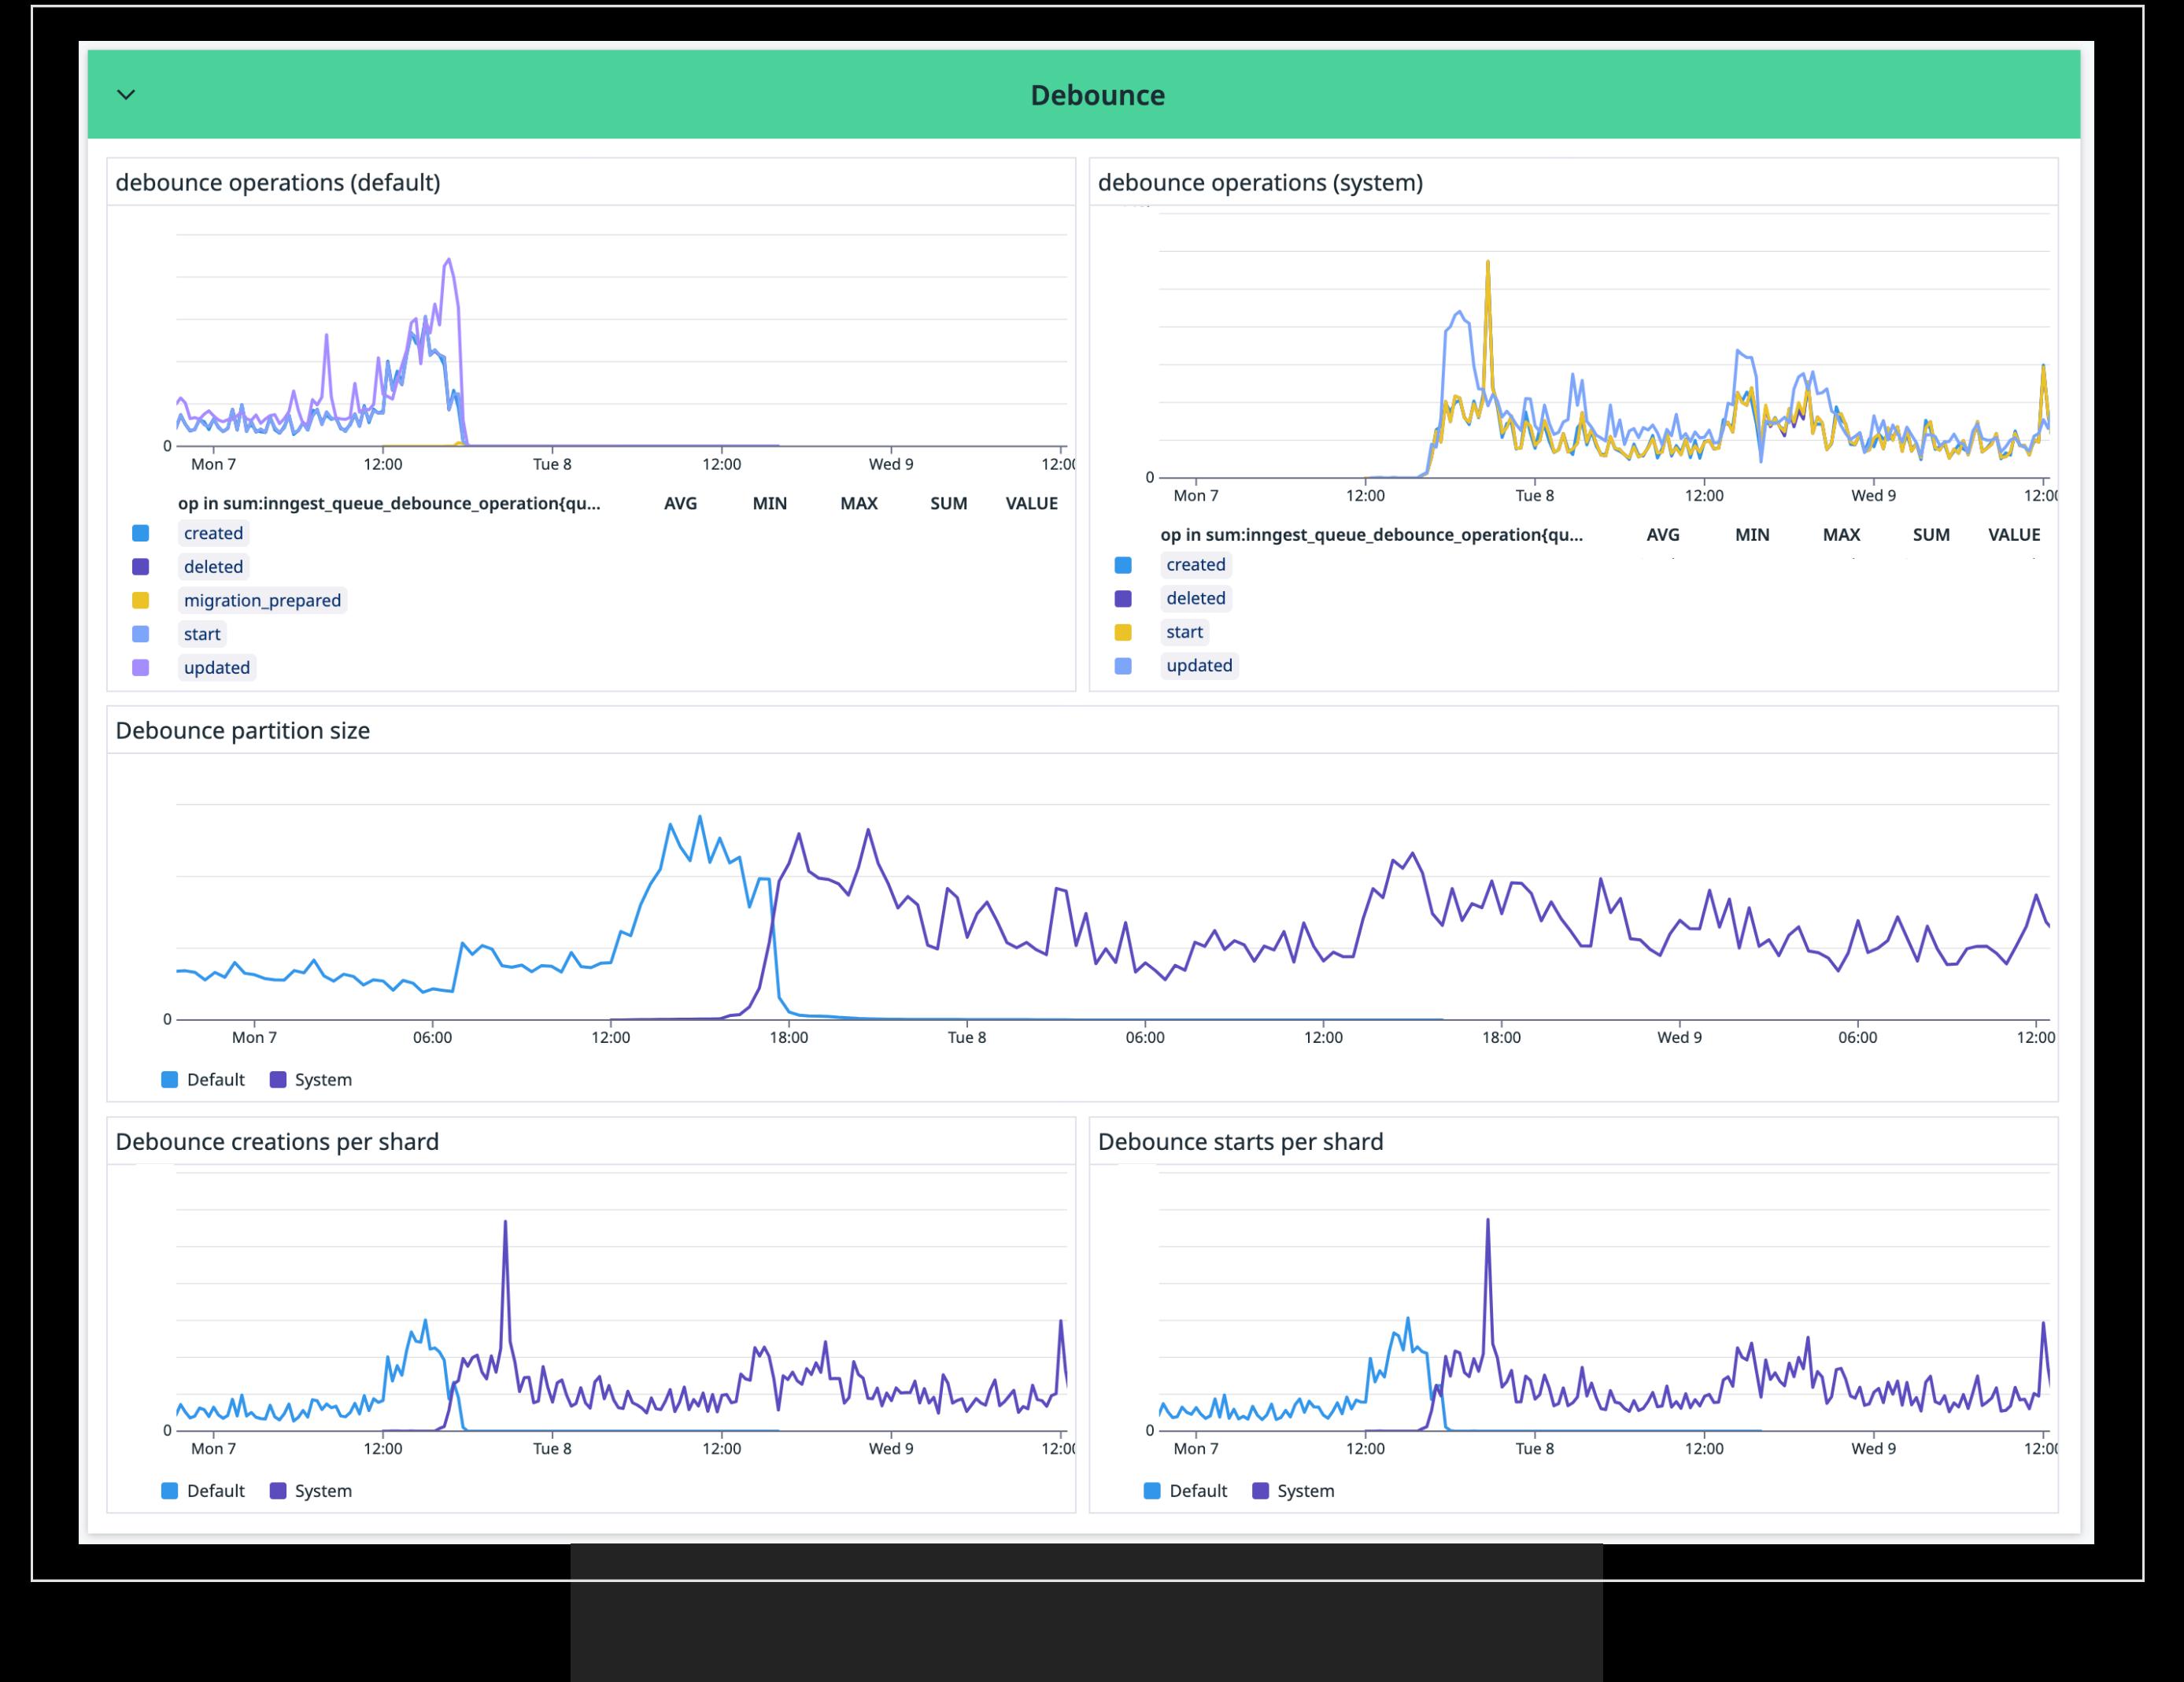


■ 14 August/25

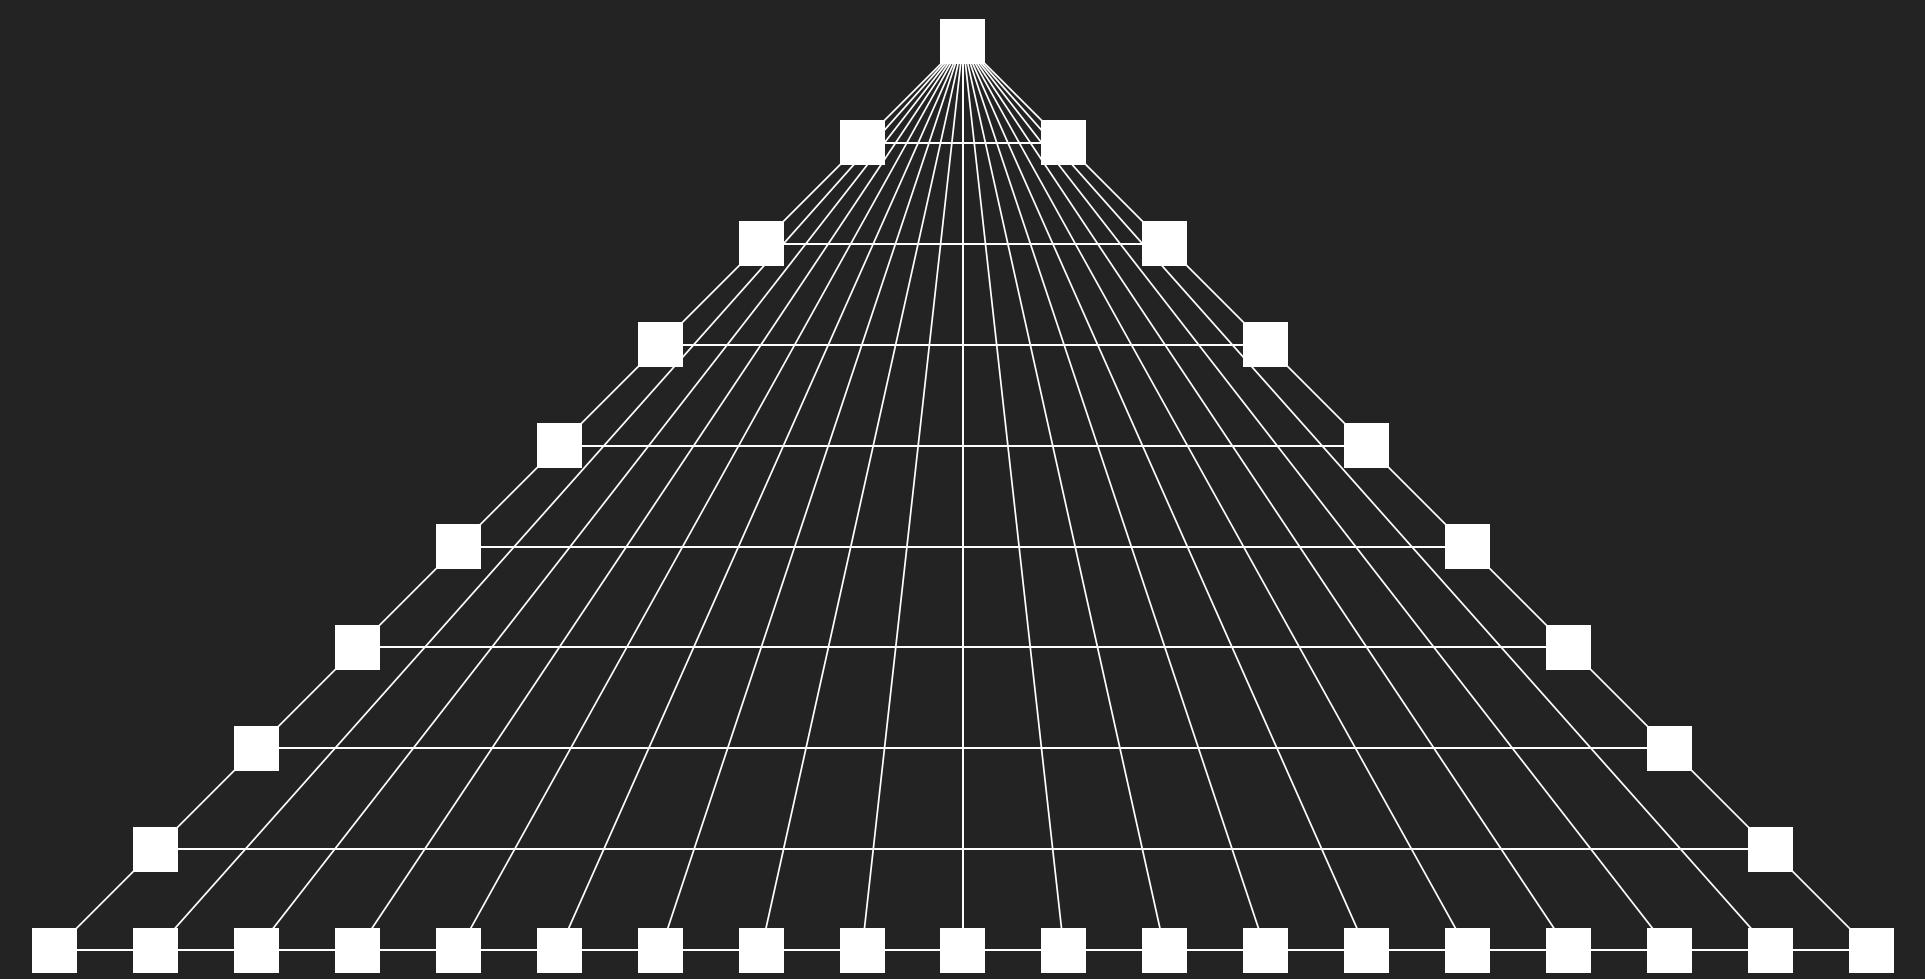
# Amazing migration







Same goal,  
(vastly)  
different  
outcomes...



## ■ Background

- 
- 
- 

[↓]

# Why should we care?

■ Why should we care?

Software  
always changes  
(migrations  
are inevitable)

[01] Dealing with intense growth

[02] Building new features

[03] Launching new products

[04] Sunsetting legacy software

## ■ Why should we care?

Migrations can speed your team, or slow you down

[01] Jump to next projects faster

[02] Sunsetting legacy software

[03] Phase out old code, technical debt

## ■ Why should we care?

Migrations can  
be fun, really!

[01] Solving a puzzle

[02] Spikes your adrenaline  
+ dopamine

[03] Mixing disciplines:  
Architecture,  
Engineering, Ops

## ■ Why should we care?

# What's in it for me?

☒ After this session you'll  
be able to:

---

[01] Design migrations for production systems, small and large

---

[02] Set up effective observability to track migration progress

---

[03] Lead migrations from start to finish

## ■ Agenda

Act 1: A primer  
on migrations (45 min)

---

[01] Migrations are systems

---

[02] Do the work upfront

---

[02] Make it easy to do the  
right thing

---

Morning break: (30 min)

Act 2: Hands-on  
exercises (45 min)

---

[01] Introduction

---

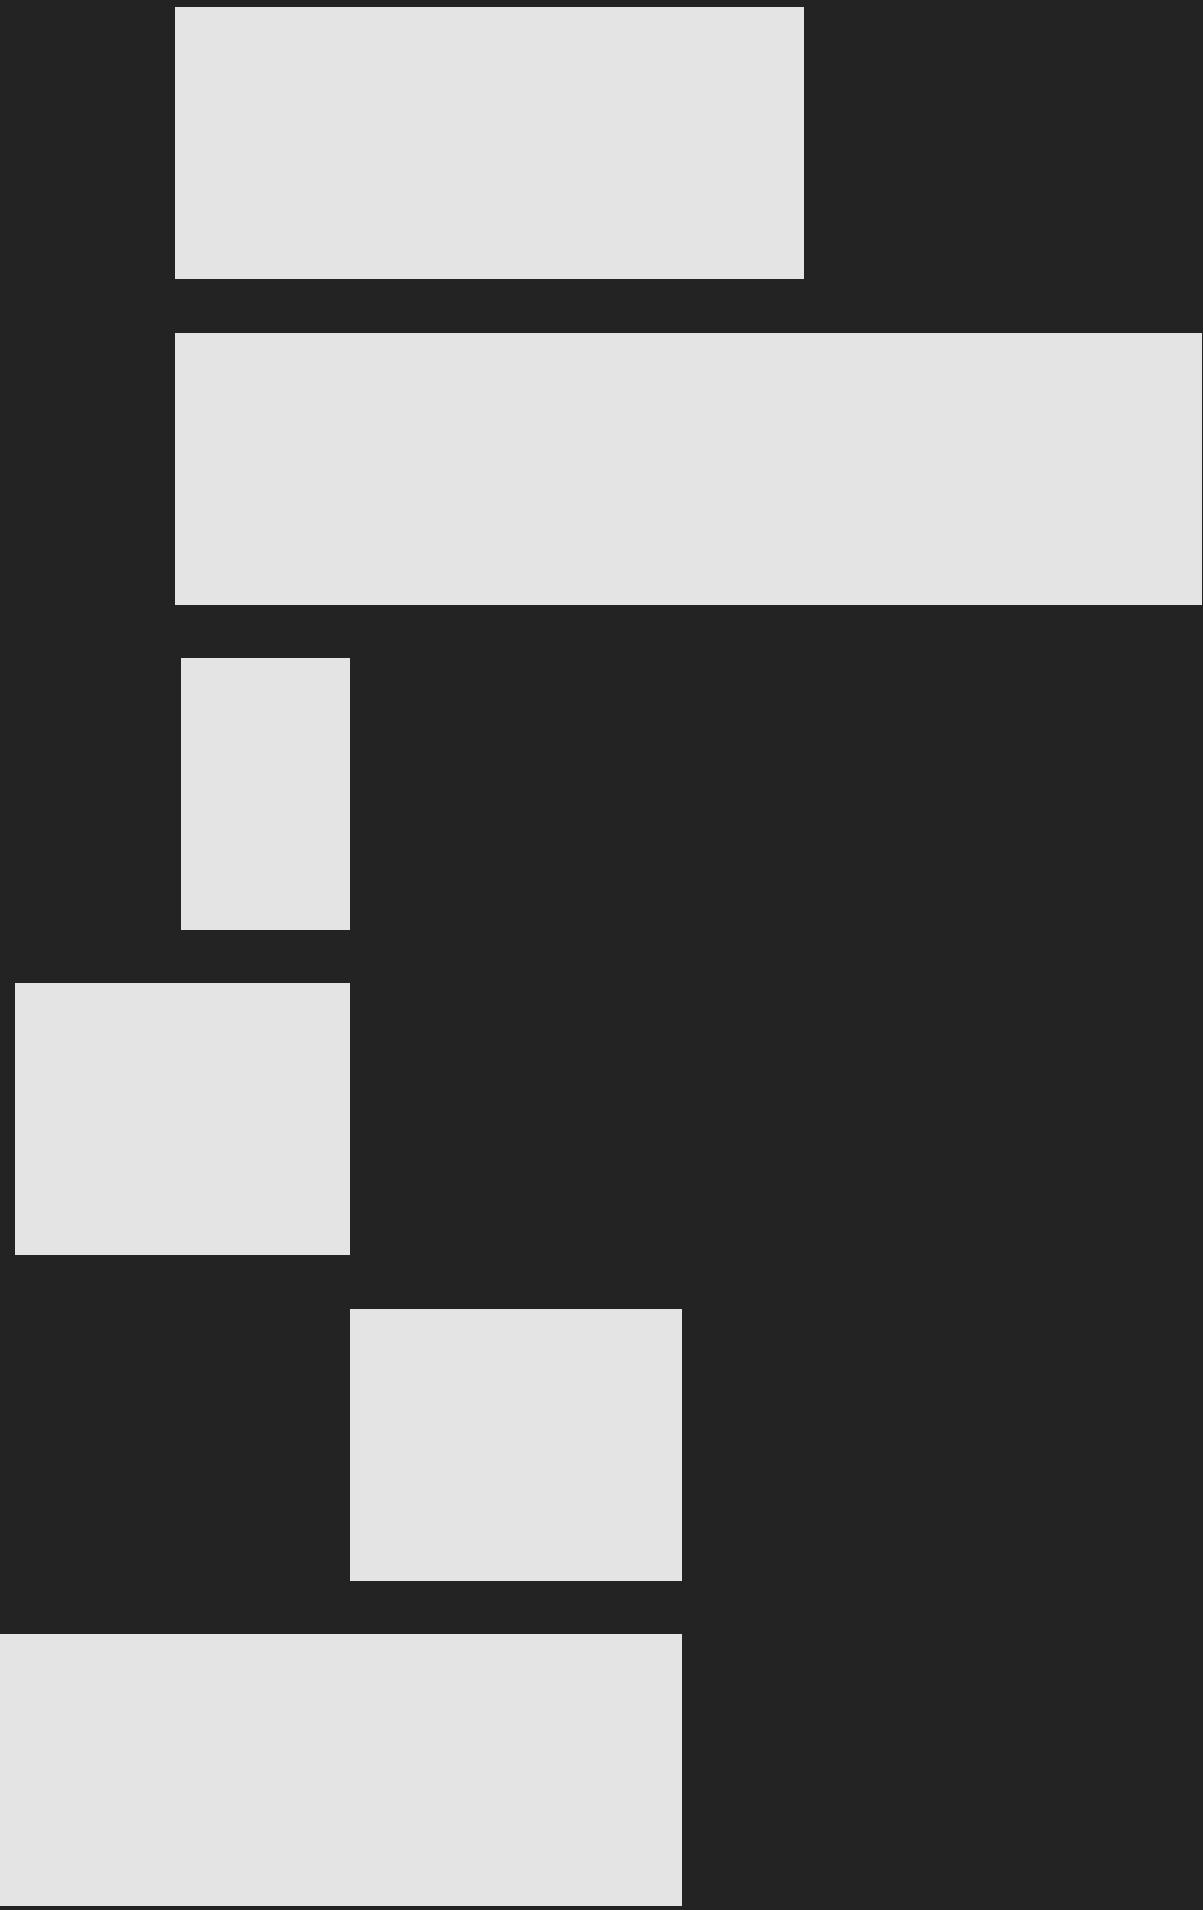
[02] Migrating data stores

---

[03] Sharding your data

---

Just ask  
questions,  
anytime



ACT →

01

## ■ DEFINITIONS

- 
- 
- 

# What's a migration?

■ THE SHORT ANSWER

A well-defined unit of change  
applied to a system

## ■ THE LONG ANSWER

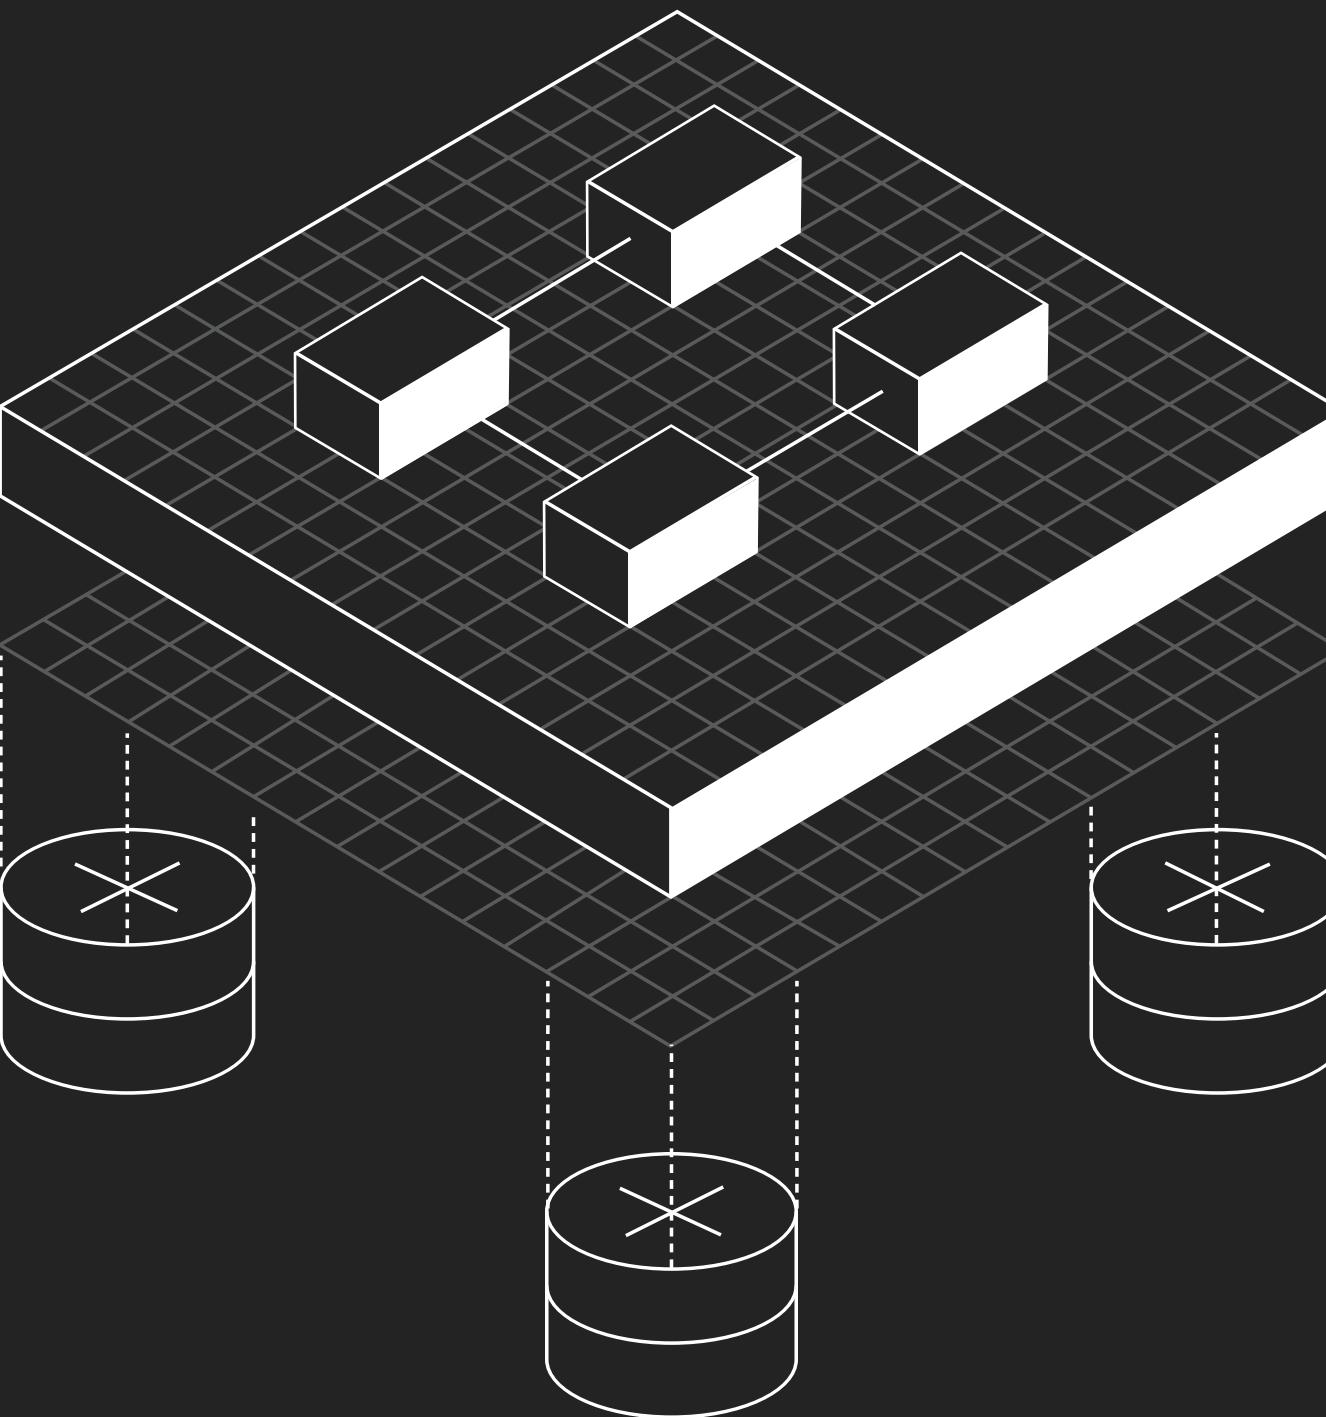
An orchestrated workflow  
made up of multiple stages with  
the goal of transitioning from  
an initial to a final state

■ The long answer

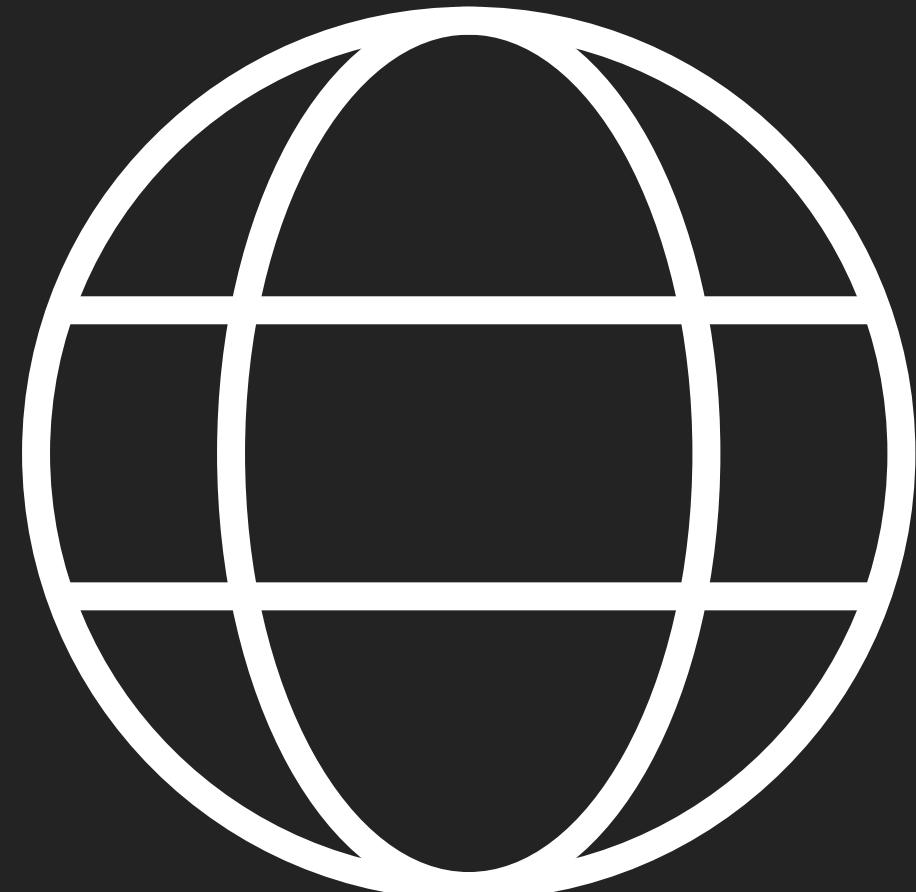
# Workflow with multiple stages, transitions

- Initial state
- Change read/ write source target
- Dual read/write
- Cutover period
- Backfill
- Final state

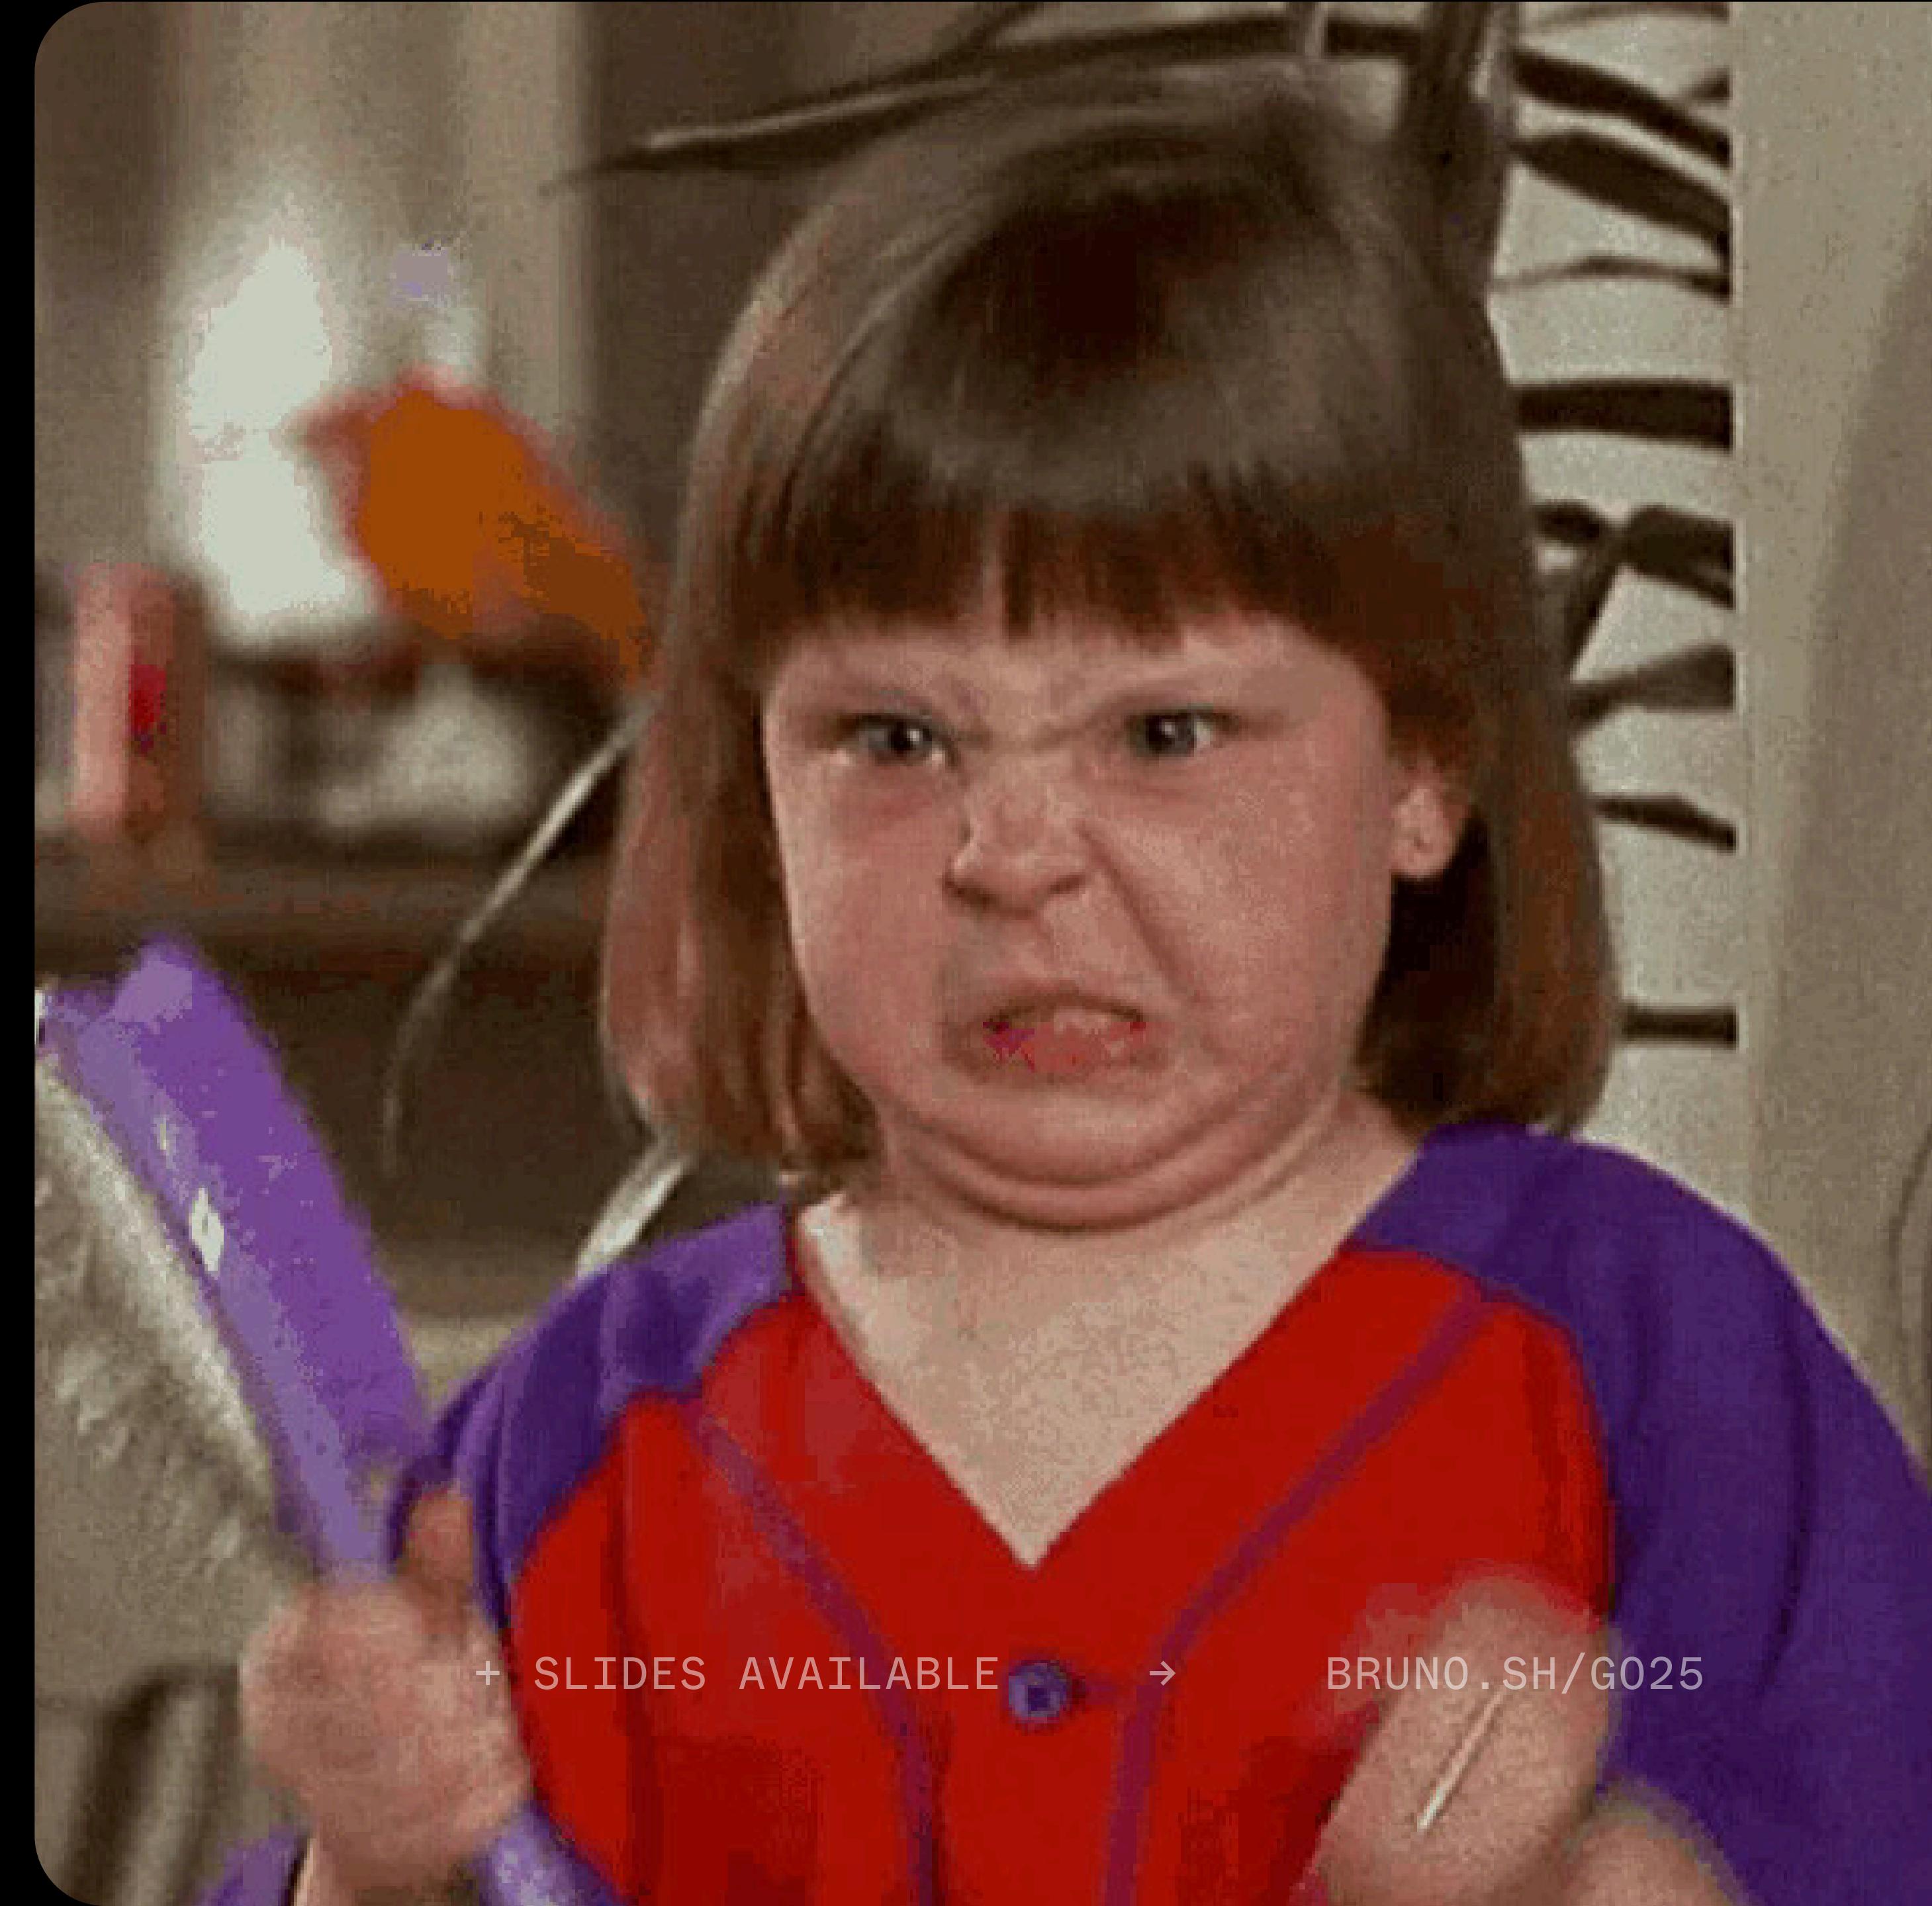
Migrations can  
apply to schema,  
data, systems,  
or a combination  
of both



Migrations can  
be online or  
offline



As a user,  
scheduled  
downtime and  
maintenance  
periods suck



■ Summary

Migrations are  
workflows  
similar to state  
machines.

[01] Migrations have  
well-defined stated  
& transition rules

[02] Migrations can affect  
schema, data & systems

[03] Migrations can be online  
& offline

# Why can't migration be more like any normal code rollout?

■ PART 1

Migrations are systems

# Three principles for migrations

- [01] Never impact a running system
- [02] Surface migration progress and controls
- [03] Always finish migrations, quickly

# Three principles for migrations

- [01] Never impact a running system
- [02] Surface migration progress and controls
- [03] Always finish migrations, quickly

# Three principles for migrations

- [01] Never impact a running system
- [02] Surface migration progress and controls
- [03] Always finish migrations, quickly

# Principles for distributed systems

- [01] Never make assumptions
- [02] Example: rolling out a simple producer/ consumer service
- [03] Everything can fail

# Principles for distributed systems

- [01] Never make assumptions
- [02] Example: rolling out a simple producer/ consumer service
- [03] Everything can fail

# Principles for distributed systems

- [01] Never make assumptions
- [02] Example: rolling out a simple producer/ consumer service
- [03] Everything can fail

# Why migrations fail

[01] Unexpected state/data  
in production

[02] Distributed systems

[03] Slower than expected

[04] Forever migrations

[05] Metastable failure

# Why migrations fail

[01] Unexpected state/data  
in production

[02] Distributed systems

[03] Slower than expected

[04] Forever migrations

[05] Metastable failure

# Why migrations fail

[01] Unexpected state/data  
in production

[02] Distributed systems

[03] Slower than expected

[04] Forever migrations

[05] Metastable failure

# Why migrations fail

[01] Unexpected state/data  
in production

[02] Distributed systems

[03] Slower than expected

[04] Forever migrations

[05] Metastable failure

# Why migrations fail

[01] Unexpected state/data  
in production

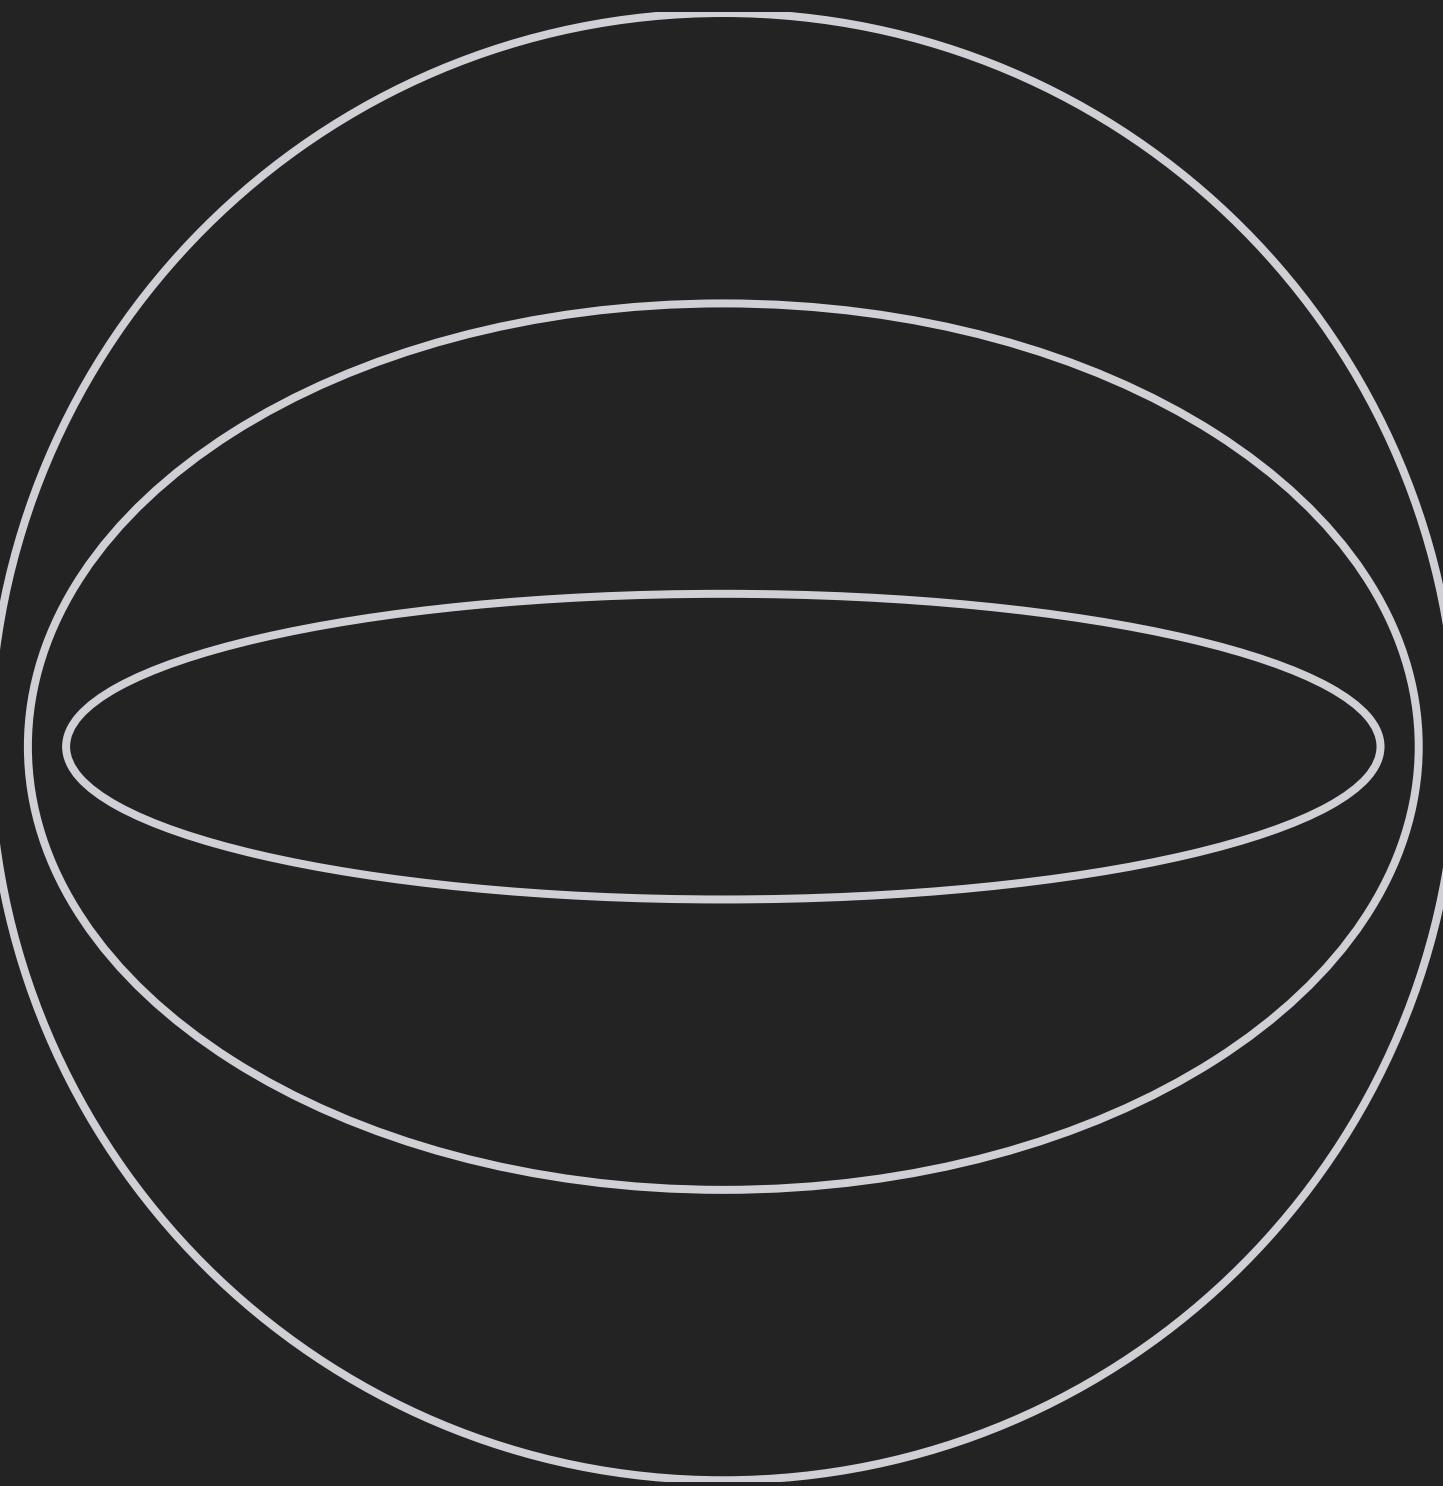
[02] Distributed systems

[03] Slower than expected

[04] Forever migrations

[05] Metastable failure

# Types of migrations (a selection)



# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks

# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks

# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks



# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks

# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks

# Service rollouts

- Intro
- Rolling releases
- Blue / green deployments
- Traffic splitting
- Shadow traffic (mirroring)
- Consistency checks

## ■ Types of migrations

# Database Schema Migrations

- Adding new tables
- Adding required columns
- Improving data access patterns with indexes, materialized views
- Backfilling data

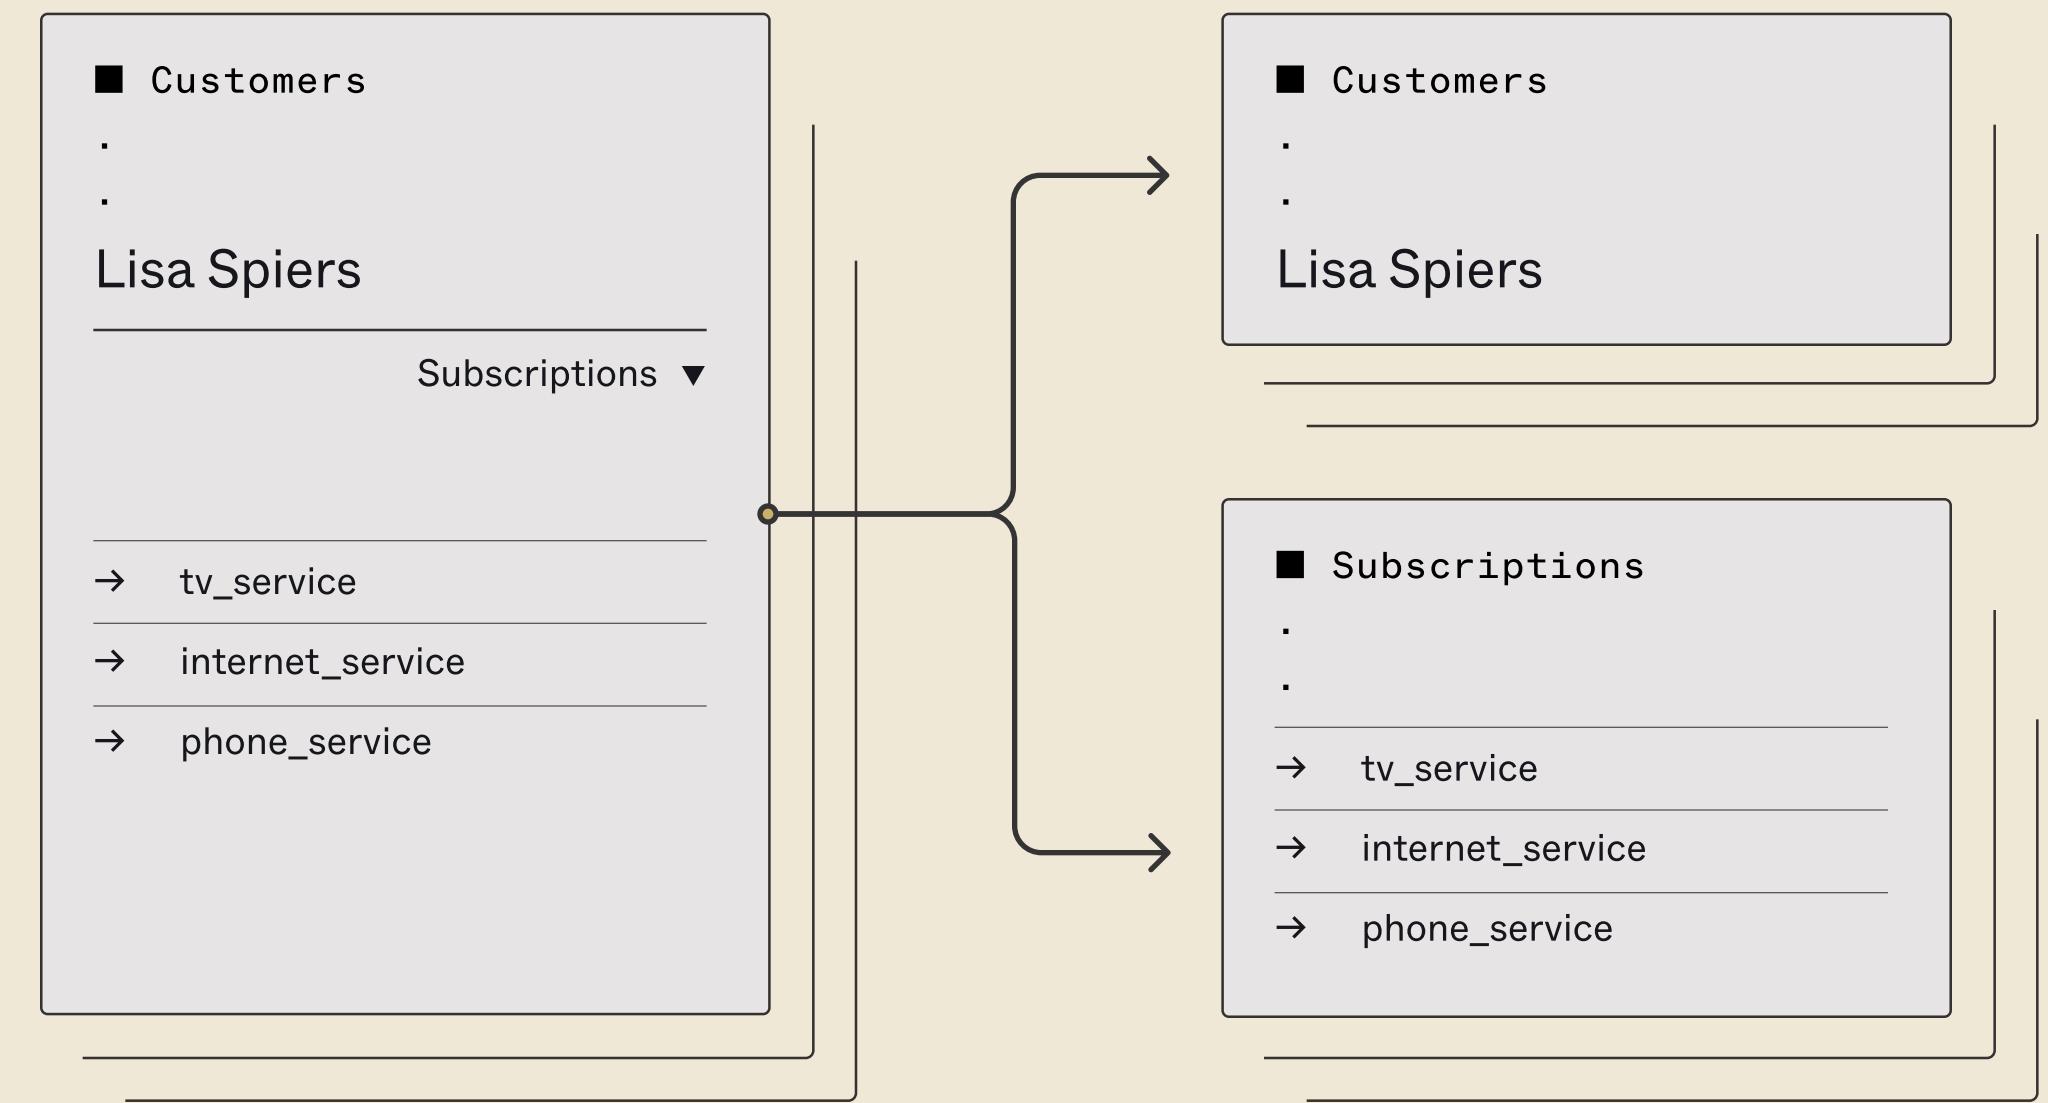
## ■ Database schema migrations

# Adding a required column (if you can't use a default value)

- Create nullable column
- Update INSERTs to always include column value for new rows
- Backfill value for existing rows
- Mark as non-nullable

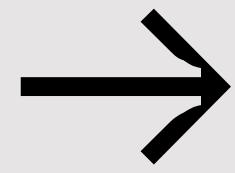
## ■ Example

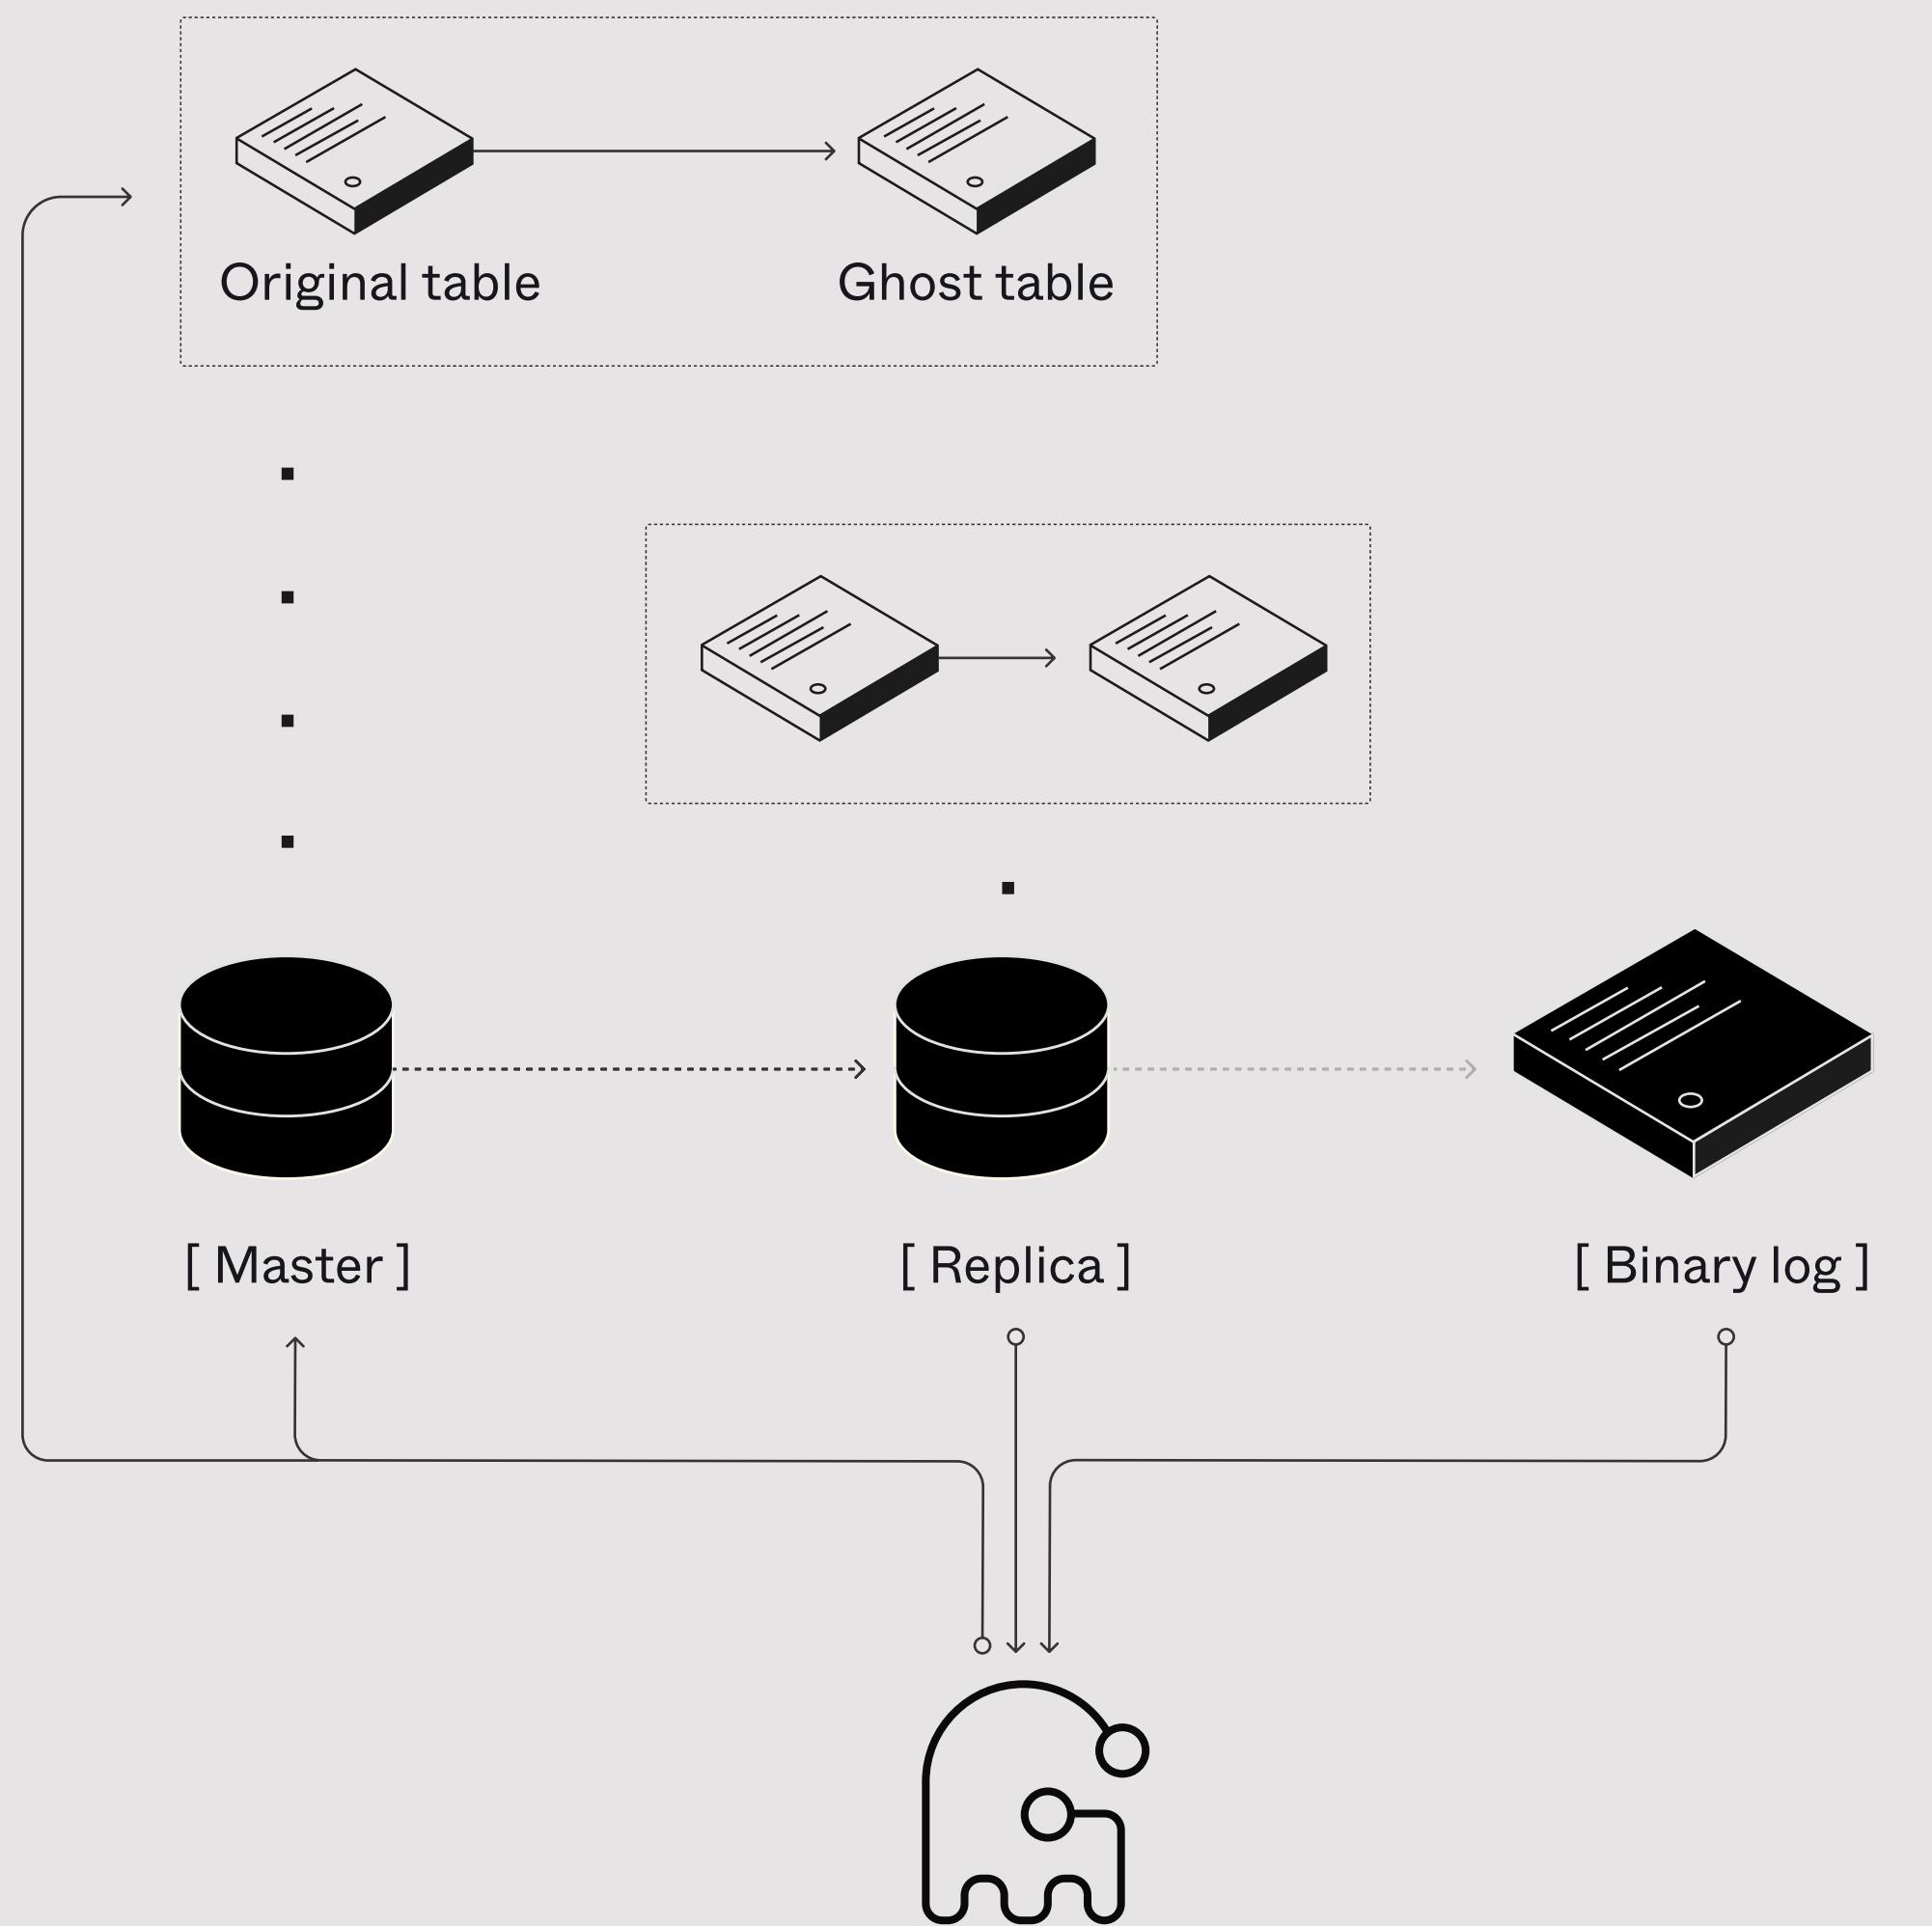
# Dual Writes (Stripe)



■ Example

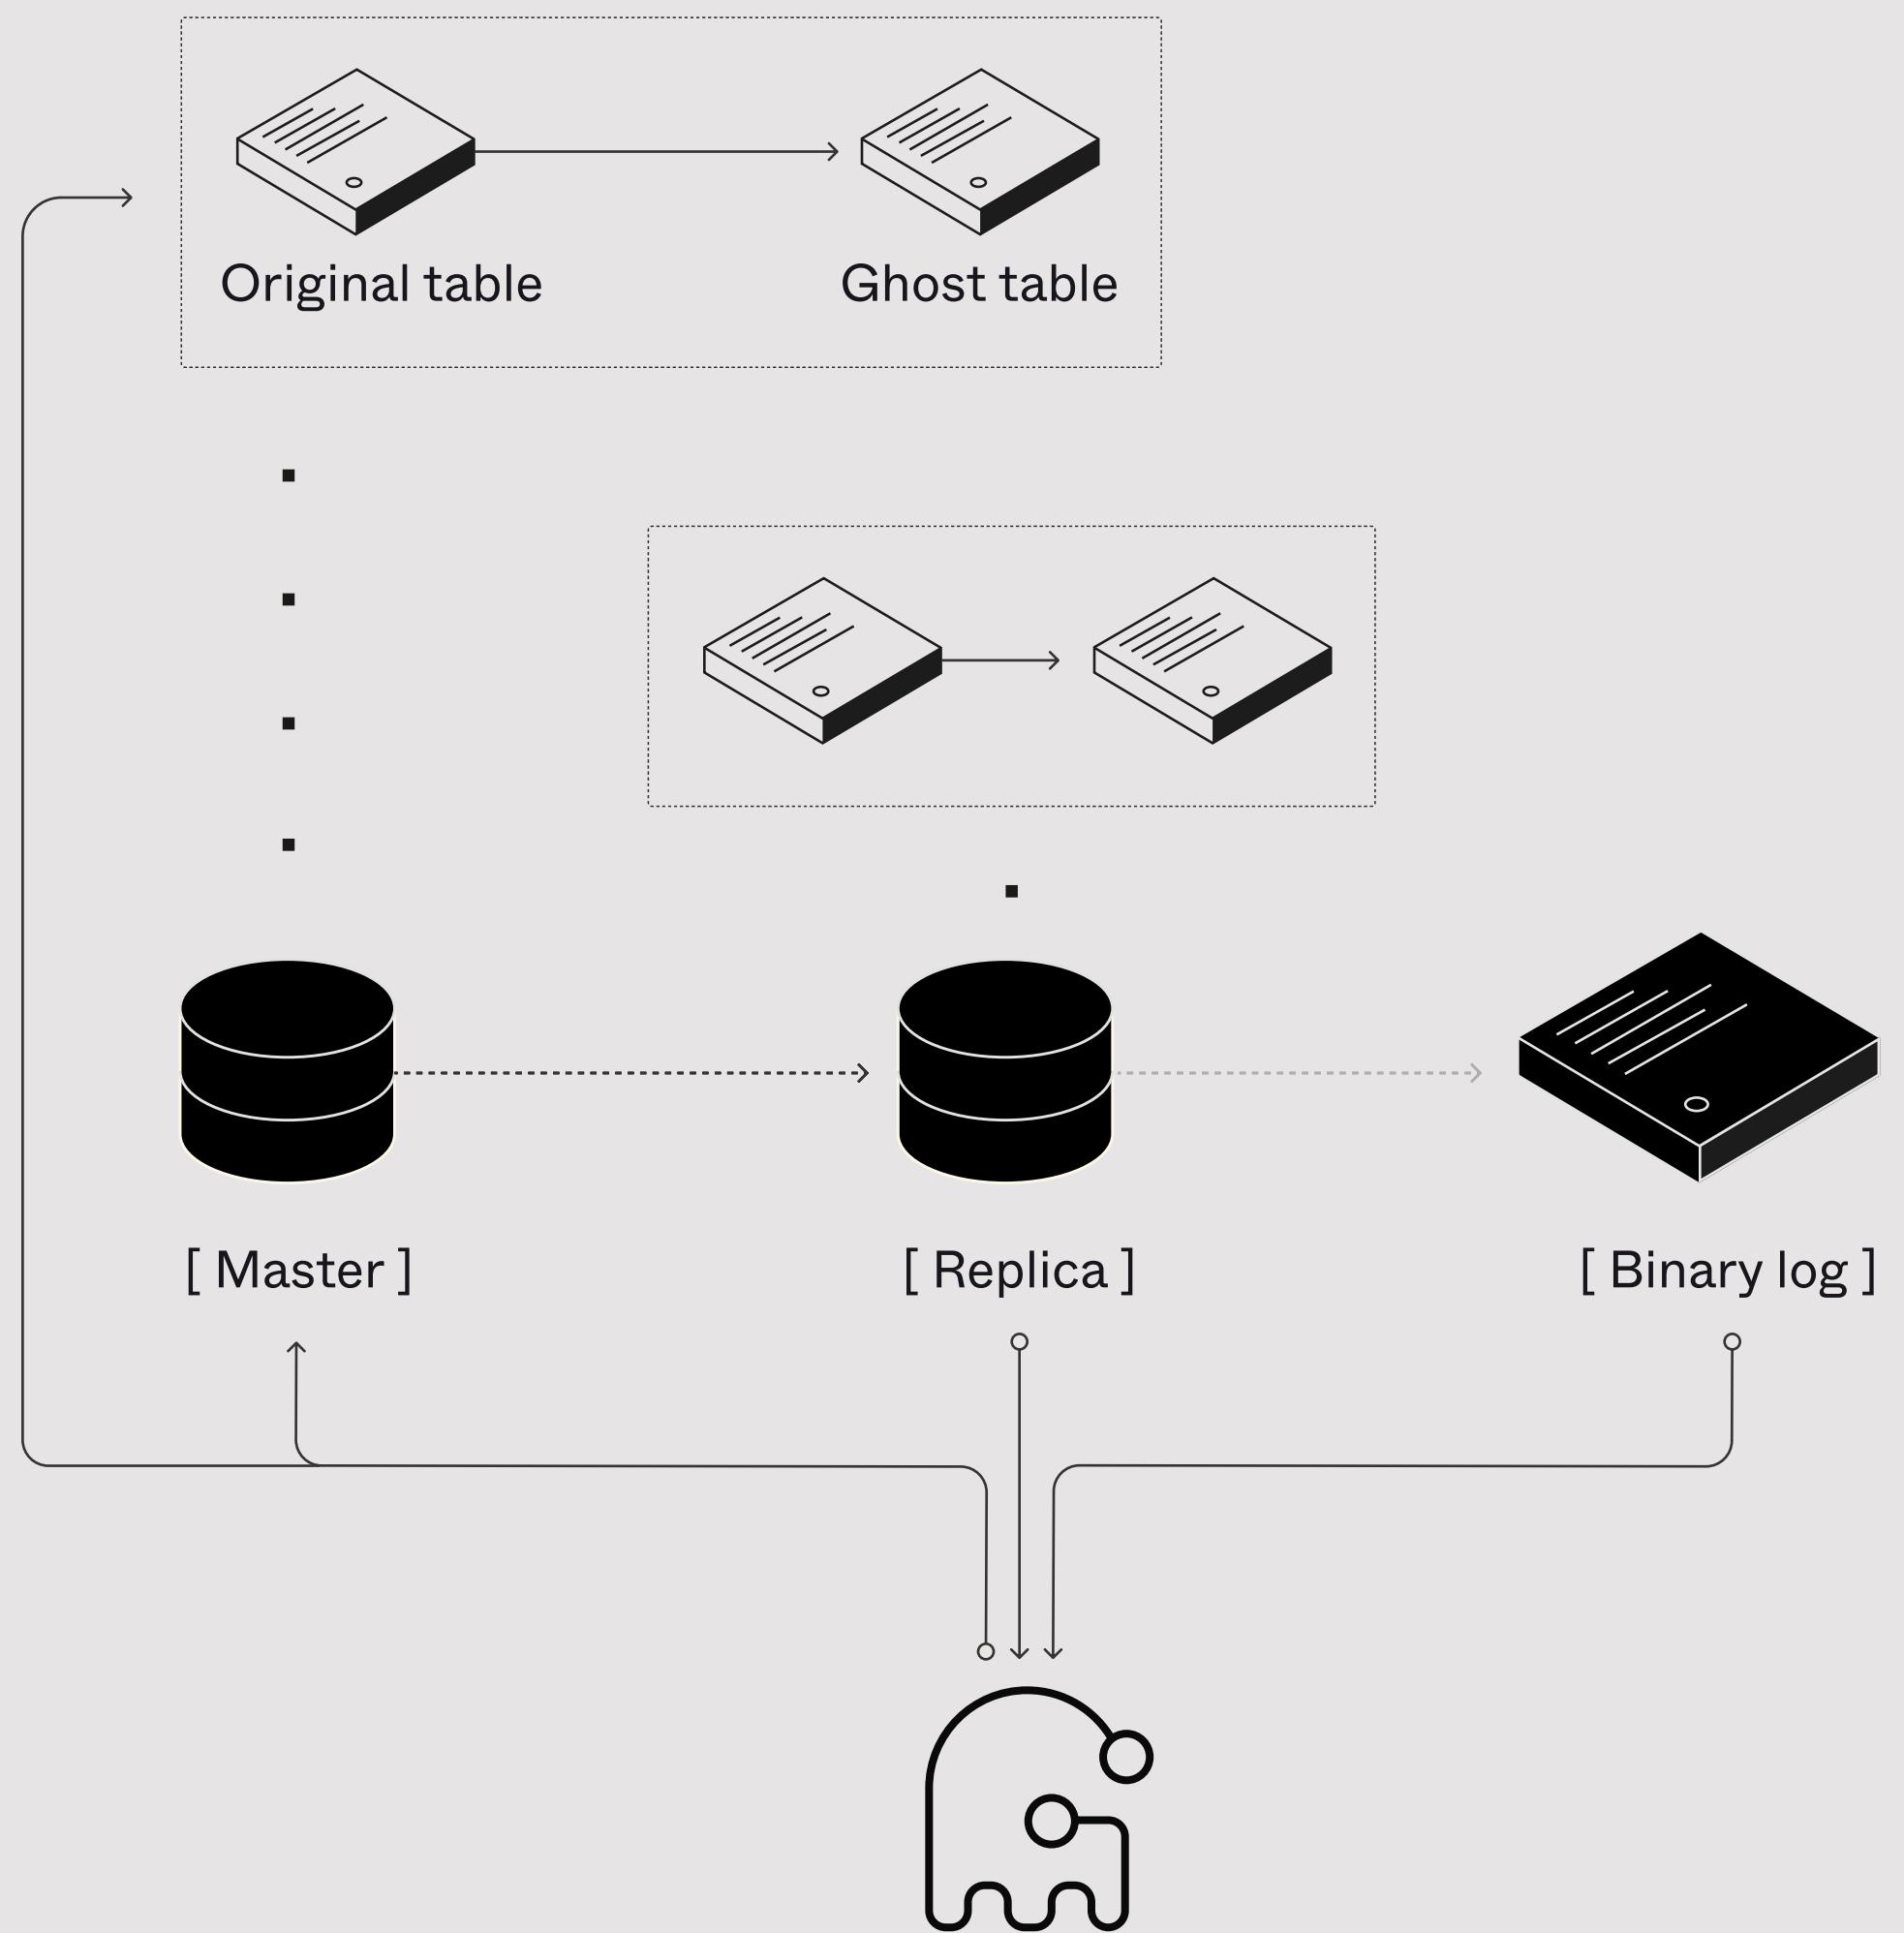
# Shadow tables (gh-ost)





## ■ gh-ost migration:

- Creates ghost table on migrated server
- Alters ghost table
- Hooks up as a MySQL replica, streams binary log events
- Interchangeably:
  - [+] Applies events on ghost table
  - [+] Copies rows from original table onto ghost table



## ■ Preferred setup:

- Connects to replica
- Inspects table structure, table dimensions on replica
- Hooks as replica onto replica
- Apply all changes on master
- Writes internal & heartbeat events onto master expects them on replica

## ■ Types of migrations

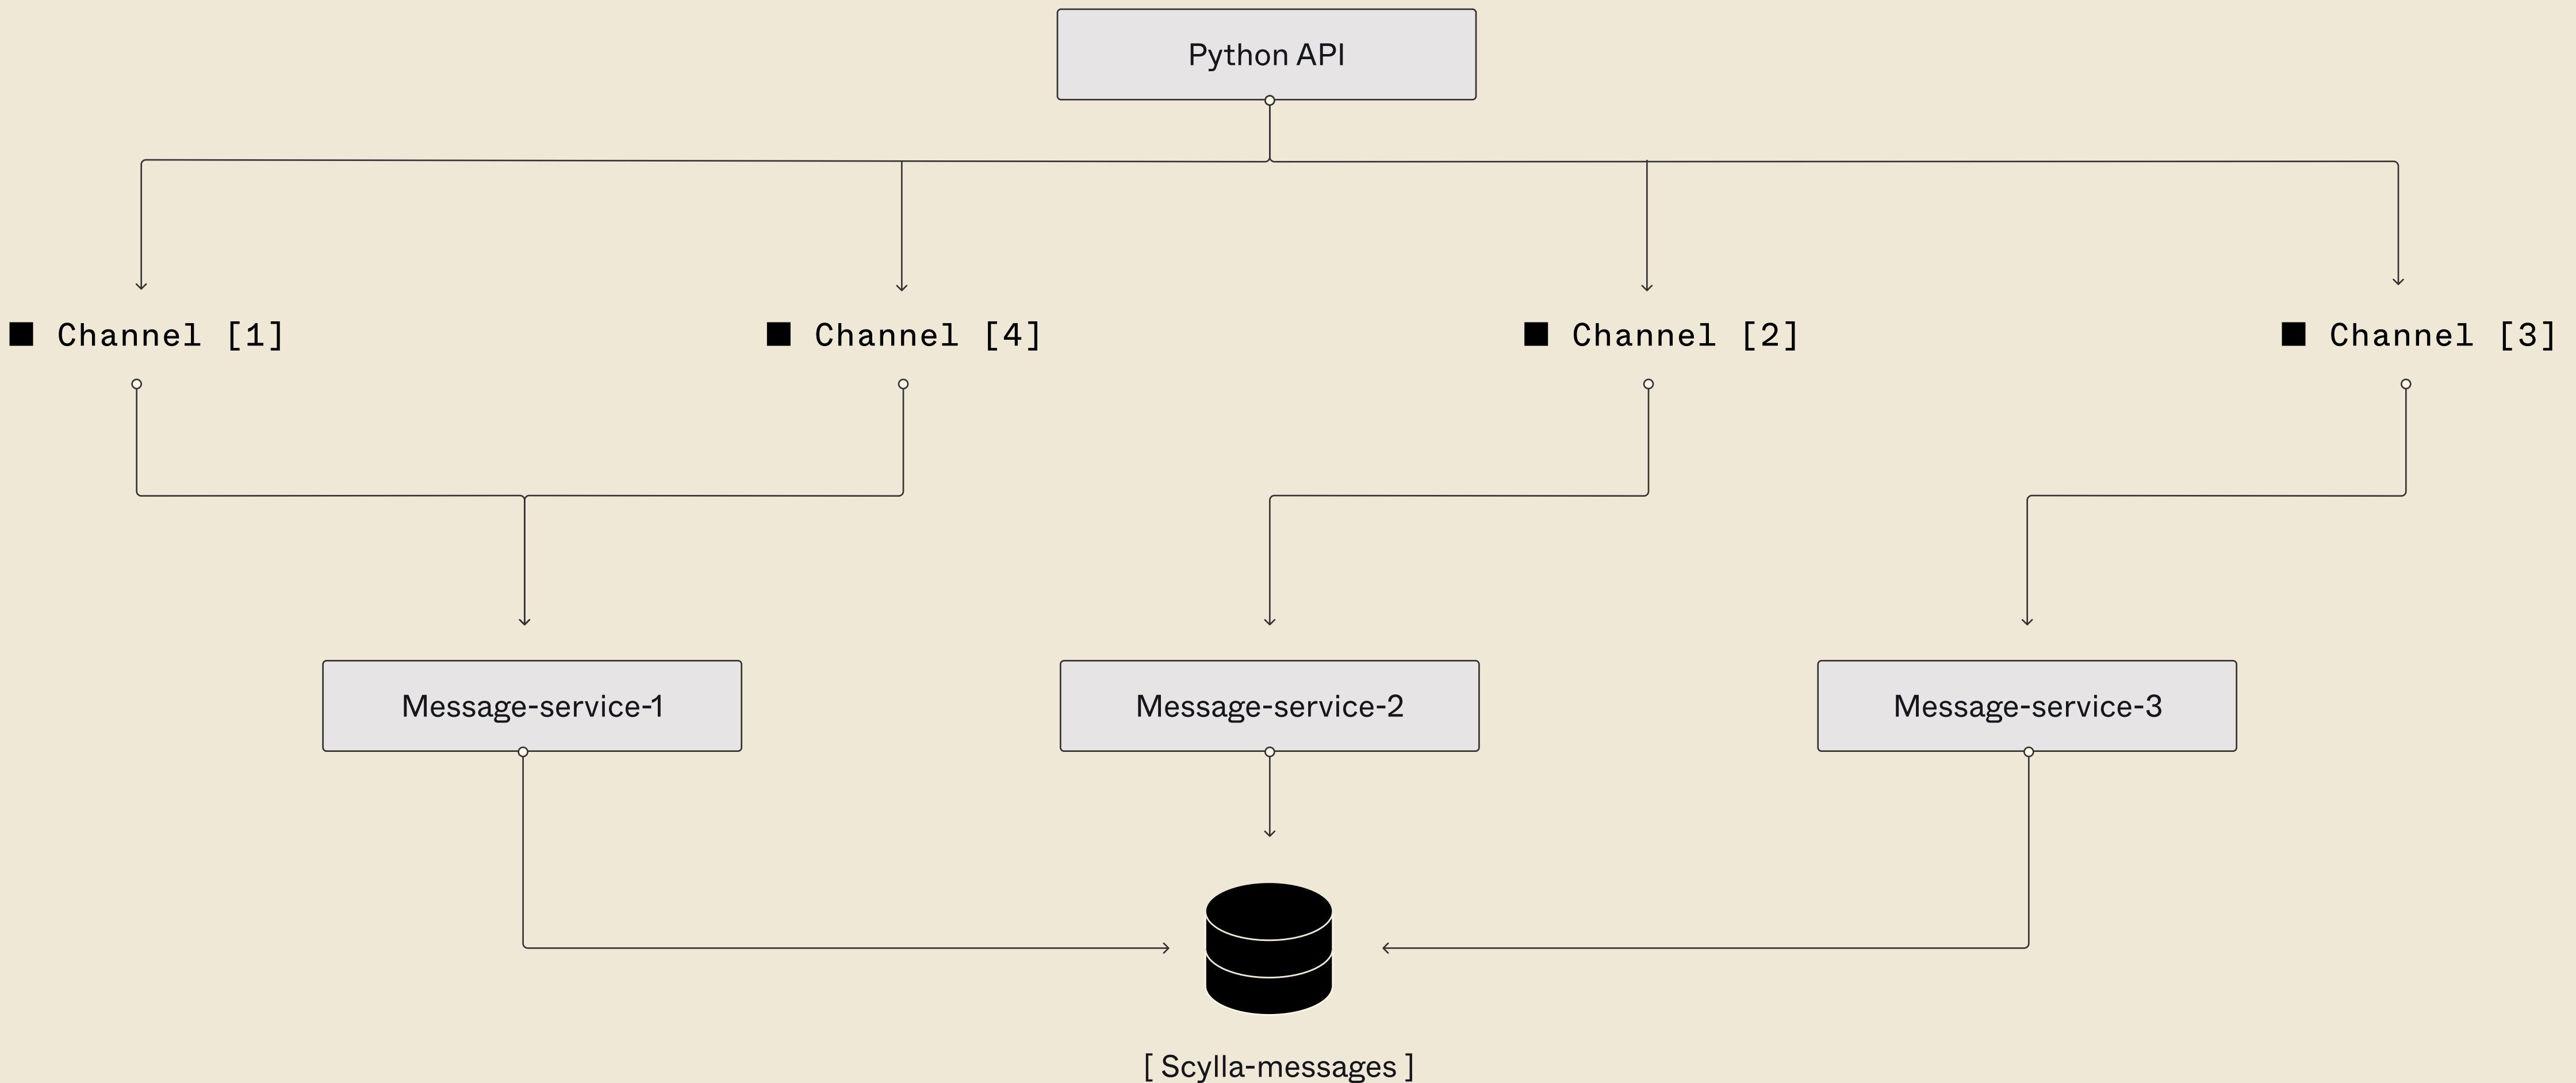
# Data Migrations

- Moving data from one data store to another
- Sharding data across multiple physical nodes
- Moving customer data from one region to another (compliance, performance)
- Isolating noisy neighbors

■ Example

# Discord migrated to ScyllaDB

- Introduced new Rust-based data (proxy) service to coalesce requests, using consistent hash-based routing
- Adding migration logic to the same data service
- Dual-wrote to both data stores during migration
- Used shadow traffic pattern to compare results between old and new clusters



■ Act 1 / Part 1

# Recap

- What are migrations
- Principles for migrations+distributed systems
- Why migrations fail
- Types of migrations

■ THESIS

Why don't we treat migrations  
the same way we plan regular  
systems?

■ Act 1 / Part 2

# Do the work upfront

- Creating a migration plan
- Never leave your migrations unattended
- Always communicate
- Make migrations a priority

- Do the work upfront

- .
- .
- .

[01]

“Time spent upfront is time  
saved during an incident”

- Do the work upfront
  - .
  - .
  - .
- [02]

# Creating a migration plan

- Which systems are involved
- What's the order of operations throughout the migration?
- Can the system be migrated gradually? What factor do you decide on?
- Can the migration be rolled back? How?

## ■ Do the work upfront

- 
- 
- 
- [03]

Never leave  
your migrations  
unattended

- Set and communicate a clear timeline (ahead of time)
- Always complete (or roll back) migrations
- Avoid “forever migrations” at all costs

## ■ Do the work upfront

- 
- 
- 
- [04]

Always  
communicate

→ Ensure all stakeholders are briefed in advance

→ Share knowledge and tools to observe systems during migration

→ Reduce everyone's stress levels!

## ■ Do the work upfront

- 
- 
- 
- [05]

Make migrations  
a priority (For  
the team leads  
among us)

- Allocate time to plan, test, and talk about migrations
- Support engineers over the entire migration lifecycle
- Do not: cut timelines, force individuals to bear the responsibility

## ■ Do the work upfront

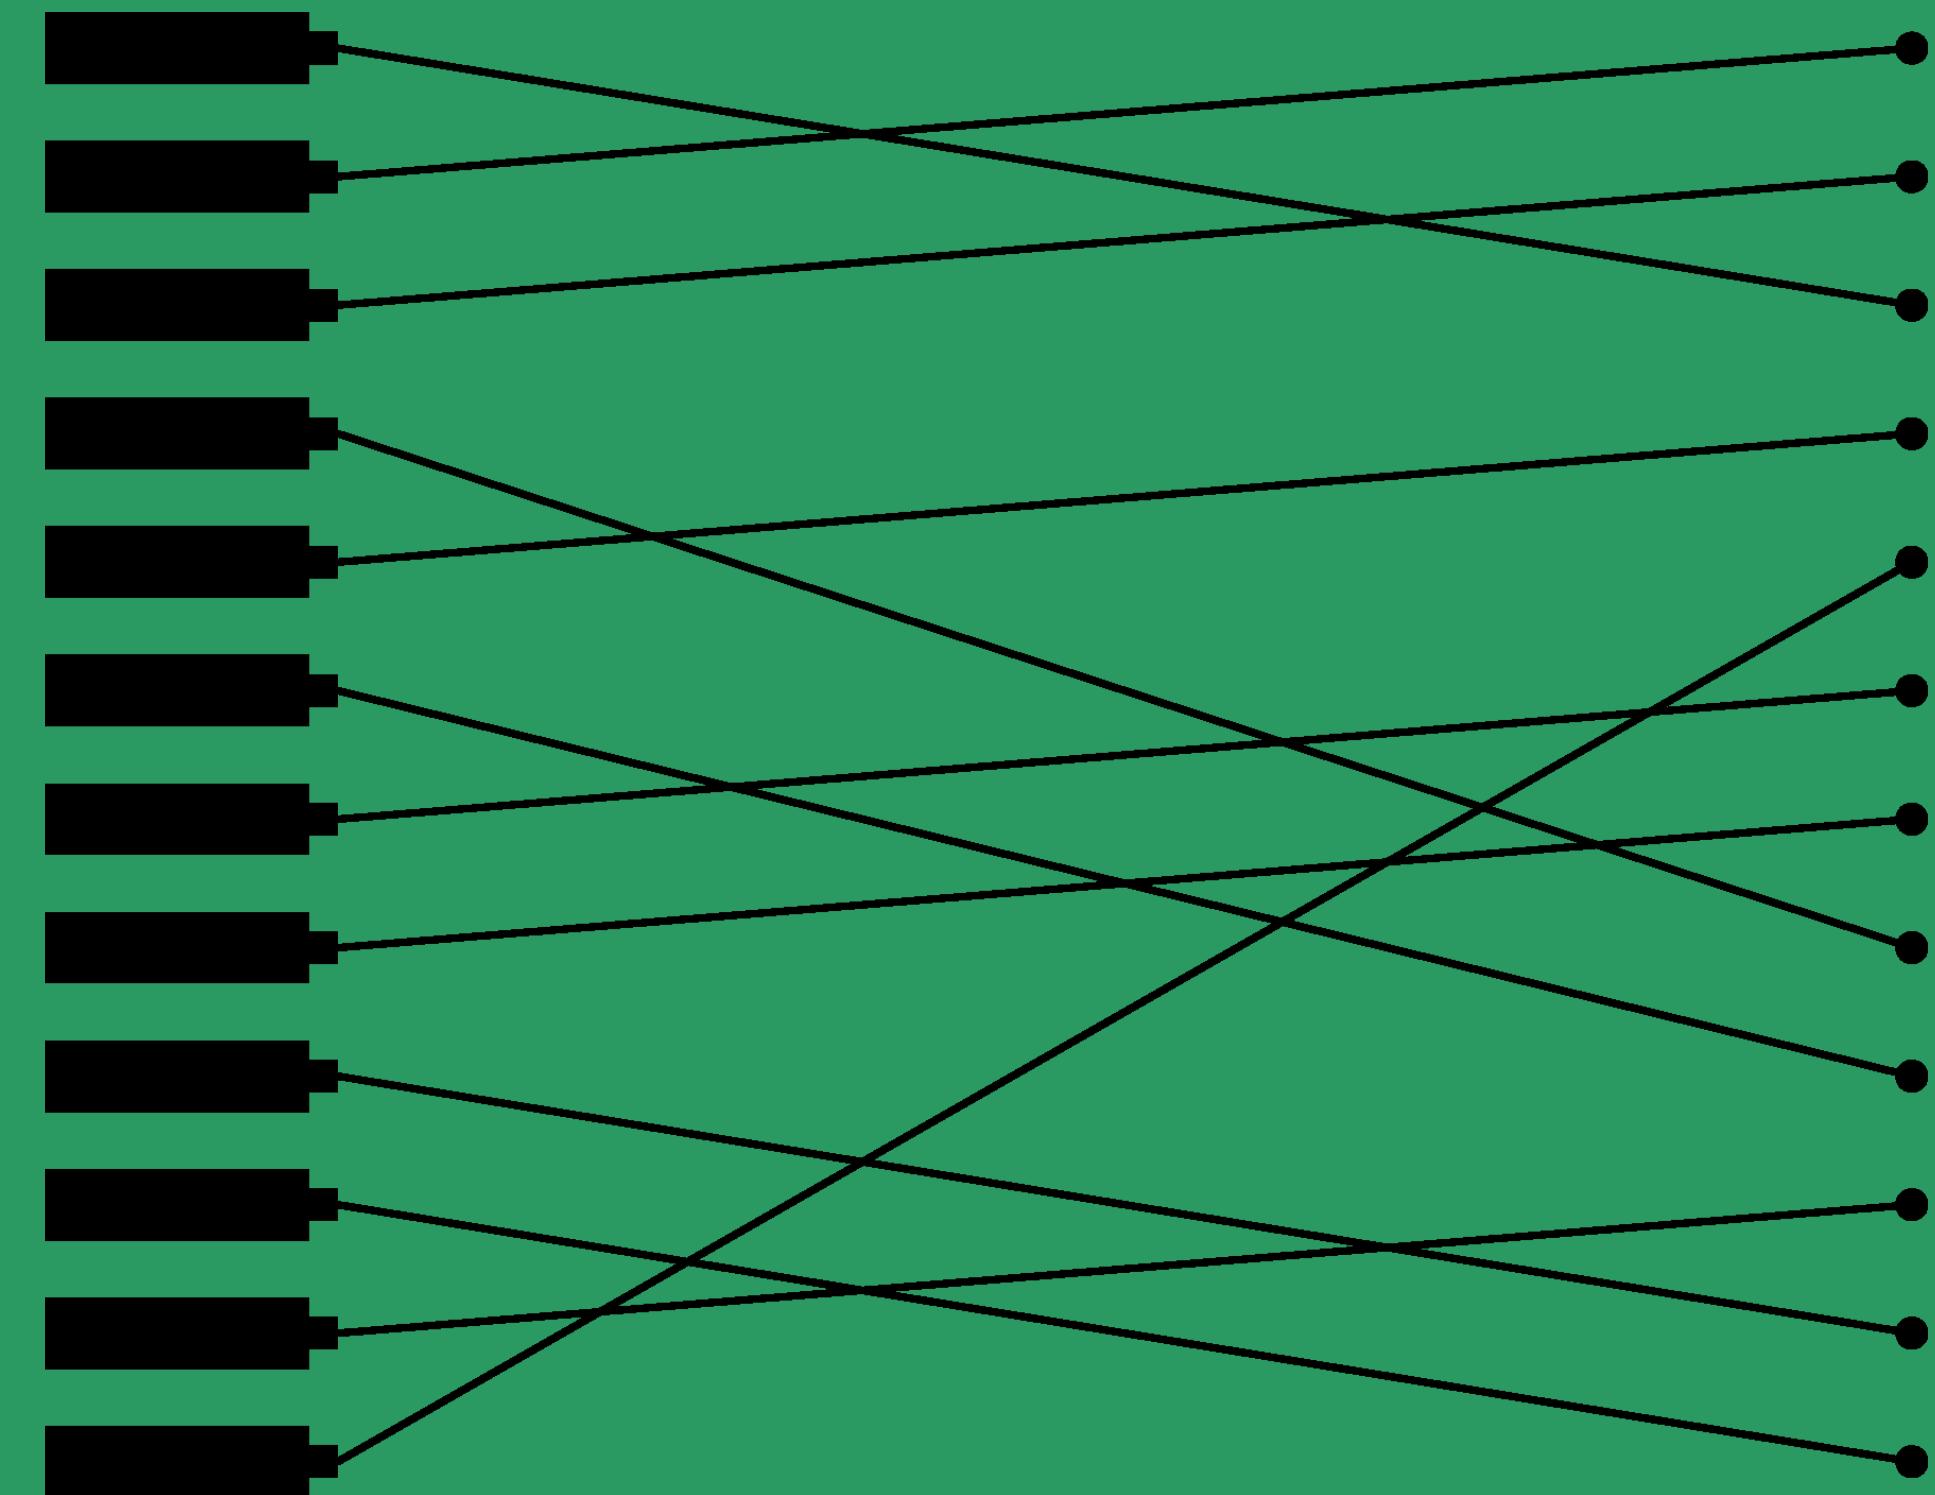
- - 
  -
- Every minute you spend planning will be paid back tenfold during execution

[06]

# Recap

■ Act 1 / Part 3

Make it easy  
(to do the right  
thing)



## ■ Make it easy

- .
- .
- .
- [01]

# Design principles (revisited)

→ Keep the scope as focused as possible

→ Find the smallest-possible unit to migrate

→ Use transactions, atomic operations

## ■ Make it easy

- .
- .
- .
- [02]

# Make your migrations visible

- Set up dashboards to follow along
- Collect fine-grained metrics to track progress
- Create alerts in case of failure

## ■ Make it easy

- .
- .
- .

[03]

Most migrations  
don't work on  
the first try

→ Test your migrations in production-like environments before rolling out

→ Aim for gradual migrations, even if it's complex to implement

→ Design migrations to be rolled back

## ■ Make it easy

- .
- .
- .

[04]

Feature flags  
are your best  
friend

→ Roll out migrations across  
a subset of customers in  
production

01 Start with internal accounts

02 Enroll free users

03 Wrap up with individual  
enterprise customers

## ■ Make it easy

- .
- .
- .

[05]

# Share knowledge early

→ Every engineer should be able  
to take over

→ Every customer-facing role  
should know the current status

## ■ Make it easy

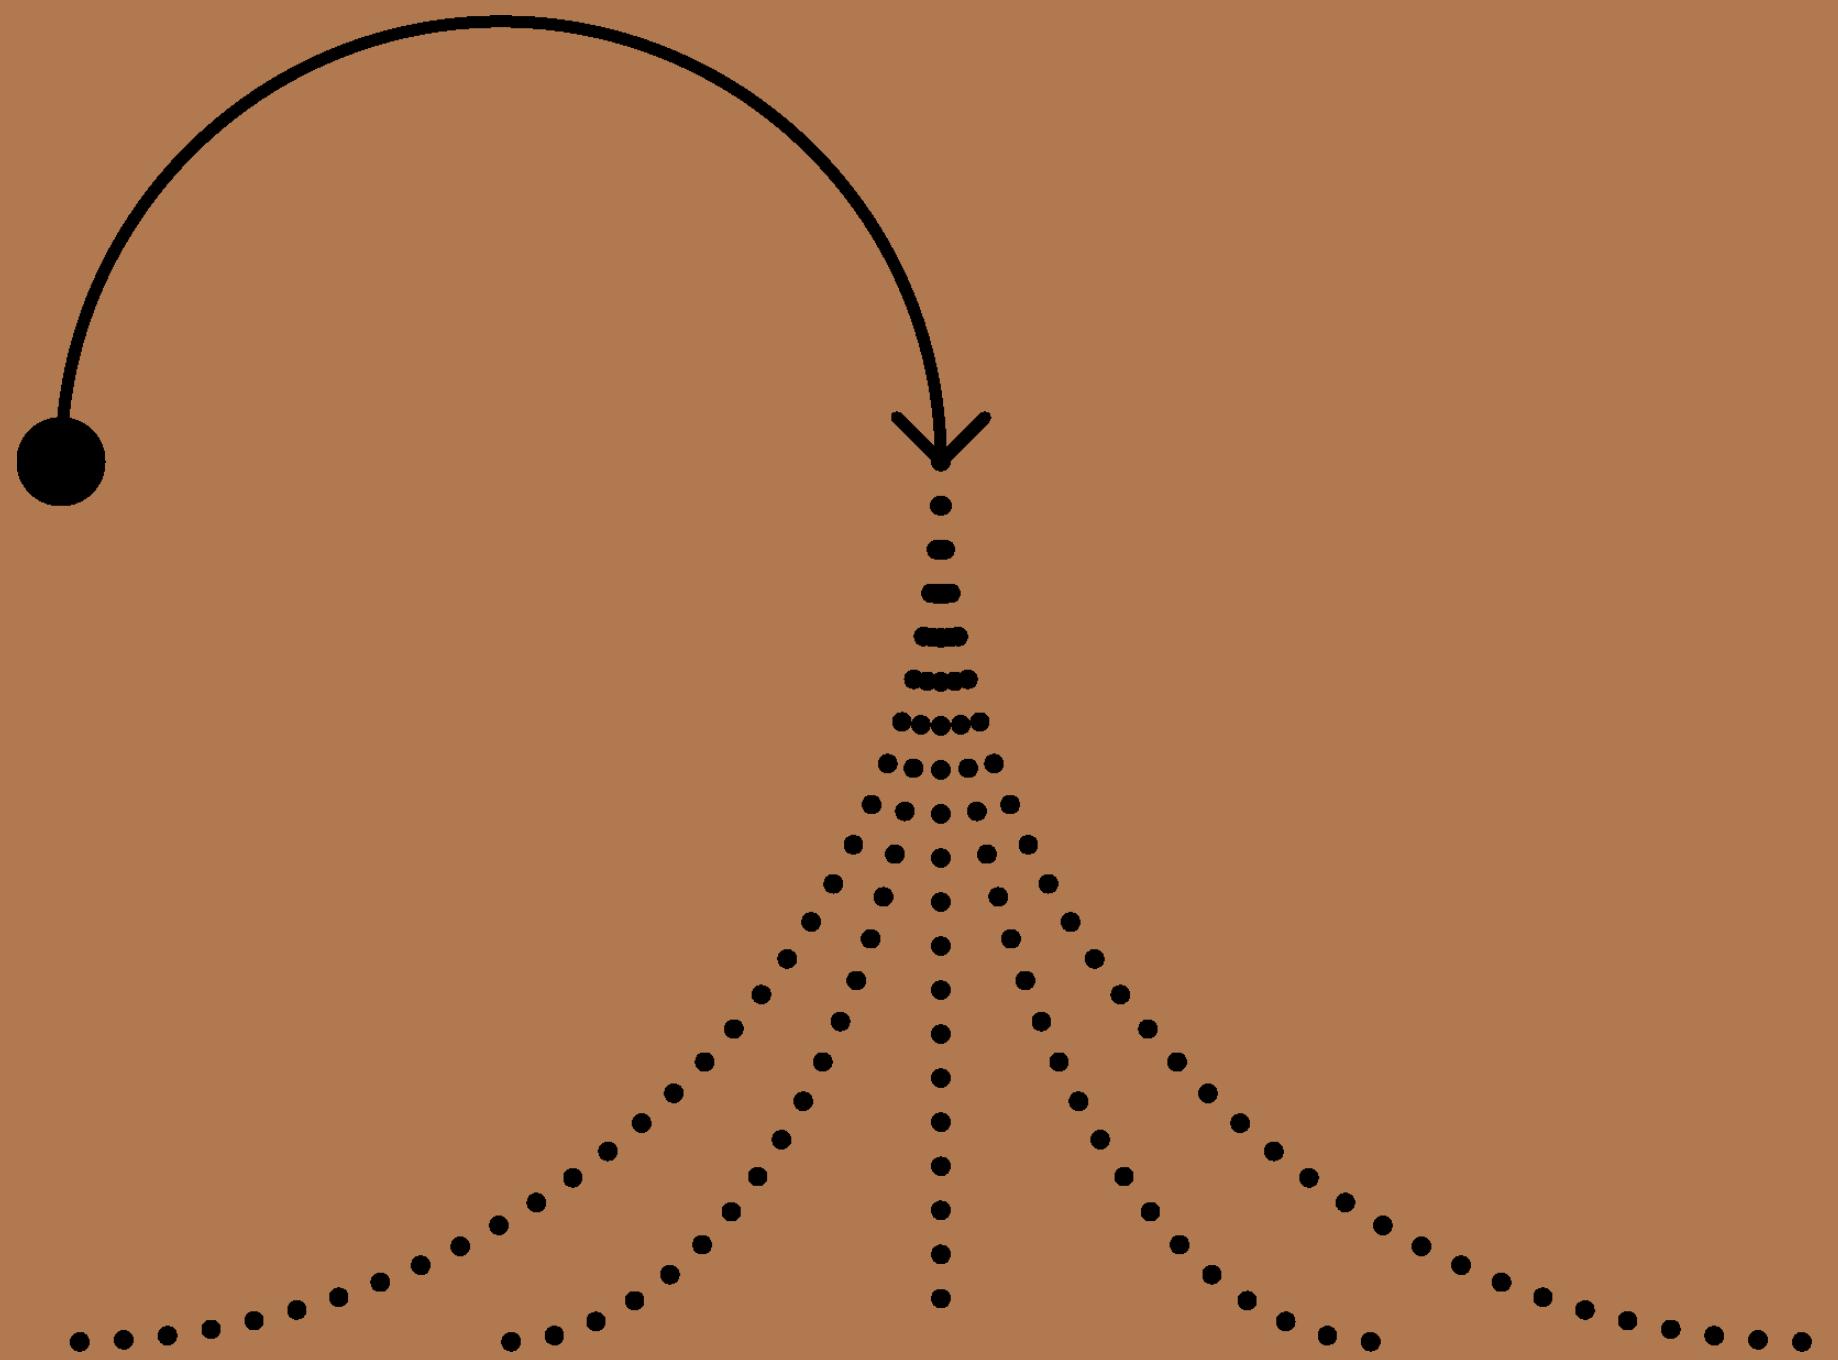
- .
- .
- .

[06]

# Recap

- Break your migrations down to the smallest units
- Track activity using metrics and visualize migration progress with dashboards
- Reduce blast radius with feature flags
- Make sure you can roll back

# Wrapping up Act 1



■ Act 1

# Recap

- Treat migrations as systems
- Do the work upfront
- Make it easy to do the right thing



.

.

.

# Before we break (What to expect next)



.

.

.

# Questions?





.

.

.

Enjoy the break,  
see you back @11:30

## ■ Hands-on exercises

- .
- .
- .
- [→]

Pilots, start  
your engines.

01 - Check out the example  
repository

→ [bruno.sh/go25](https://bruno.sh/go25)

02 - Clone the hands-on exercise  
code:

```
$ git clone github.com/  
brunoscheufler/gopherconuk25
```

03 - Start the app

→ `$ go run . --cli --gen`

ACT →

02

■ Show of hands

·

·

·

[01]

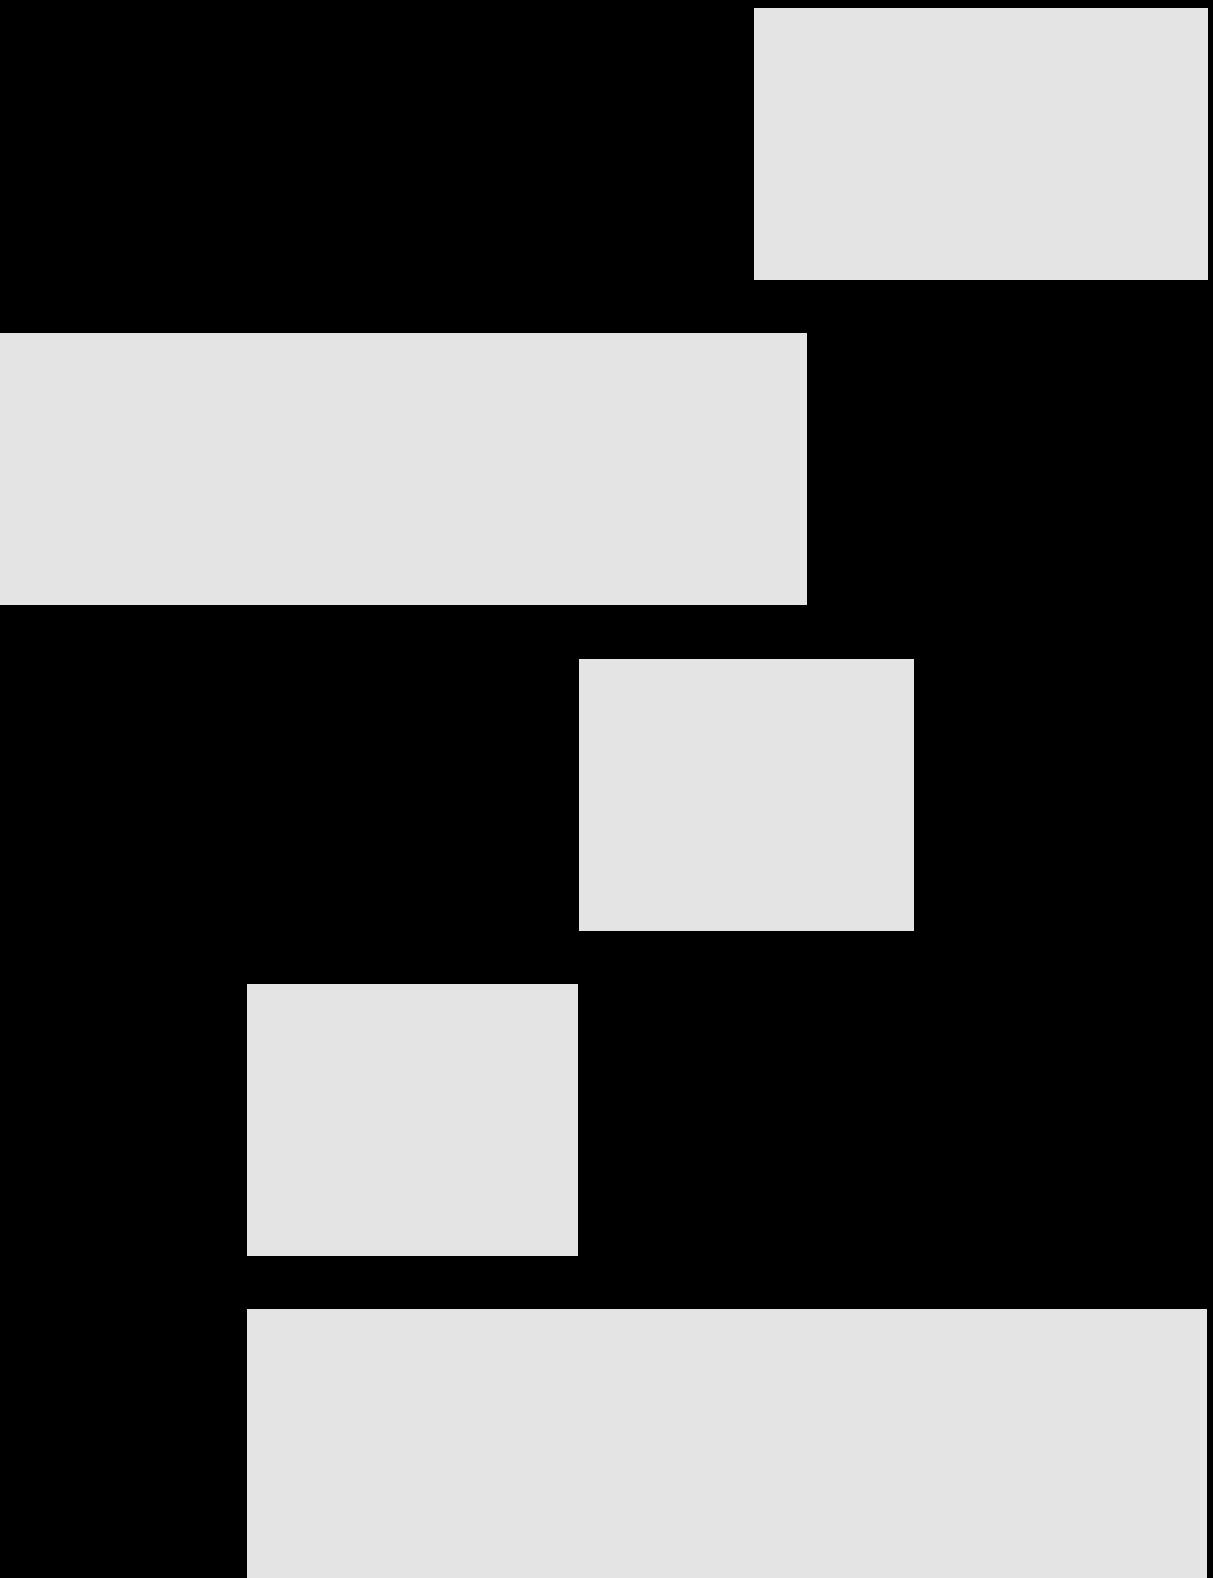
└[::/]└[ :: ]└[ / :: ]

# How much time have you spent writing Go?

# Setting up

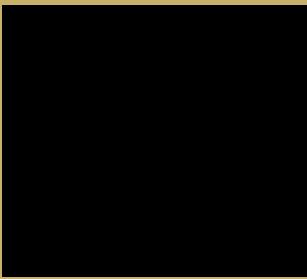
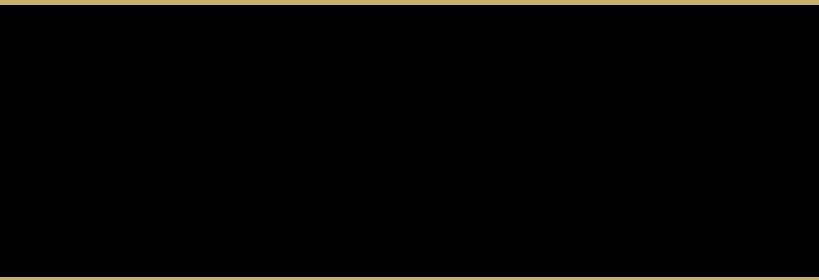
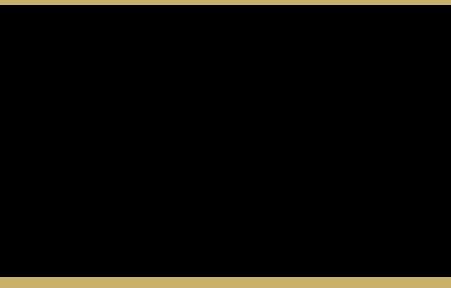


Please help  
each other,  
work together



# Our example app:

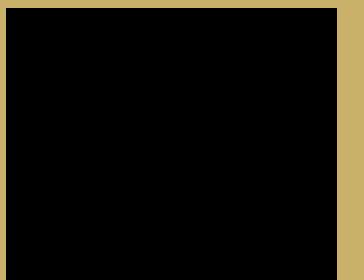
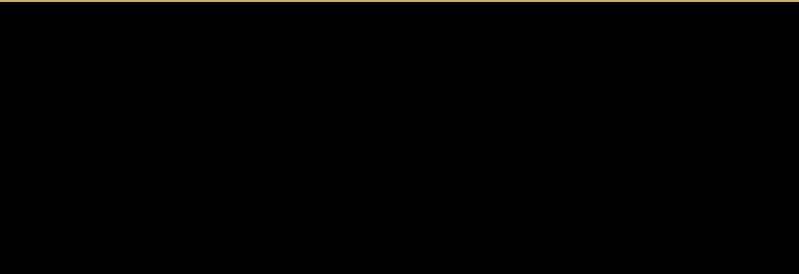
## Meet Notely



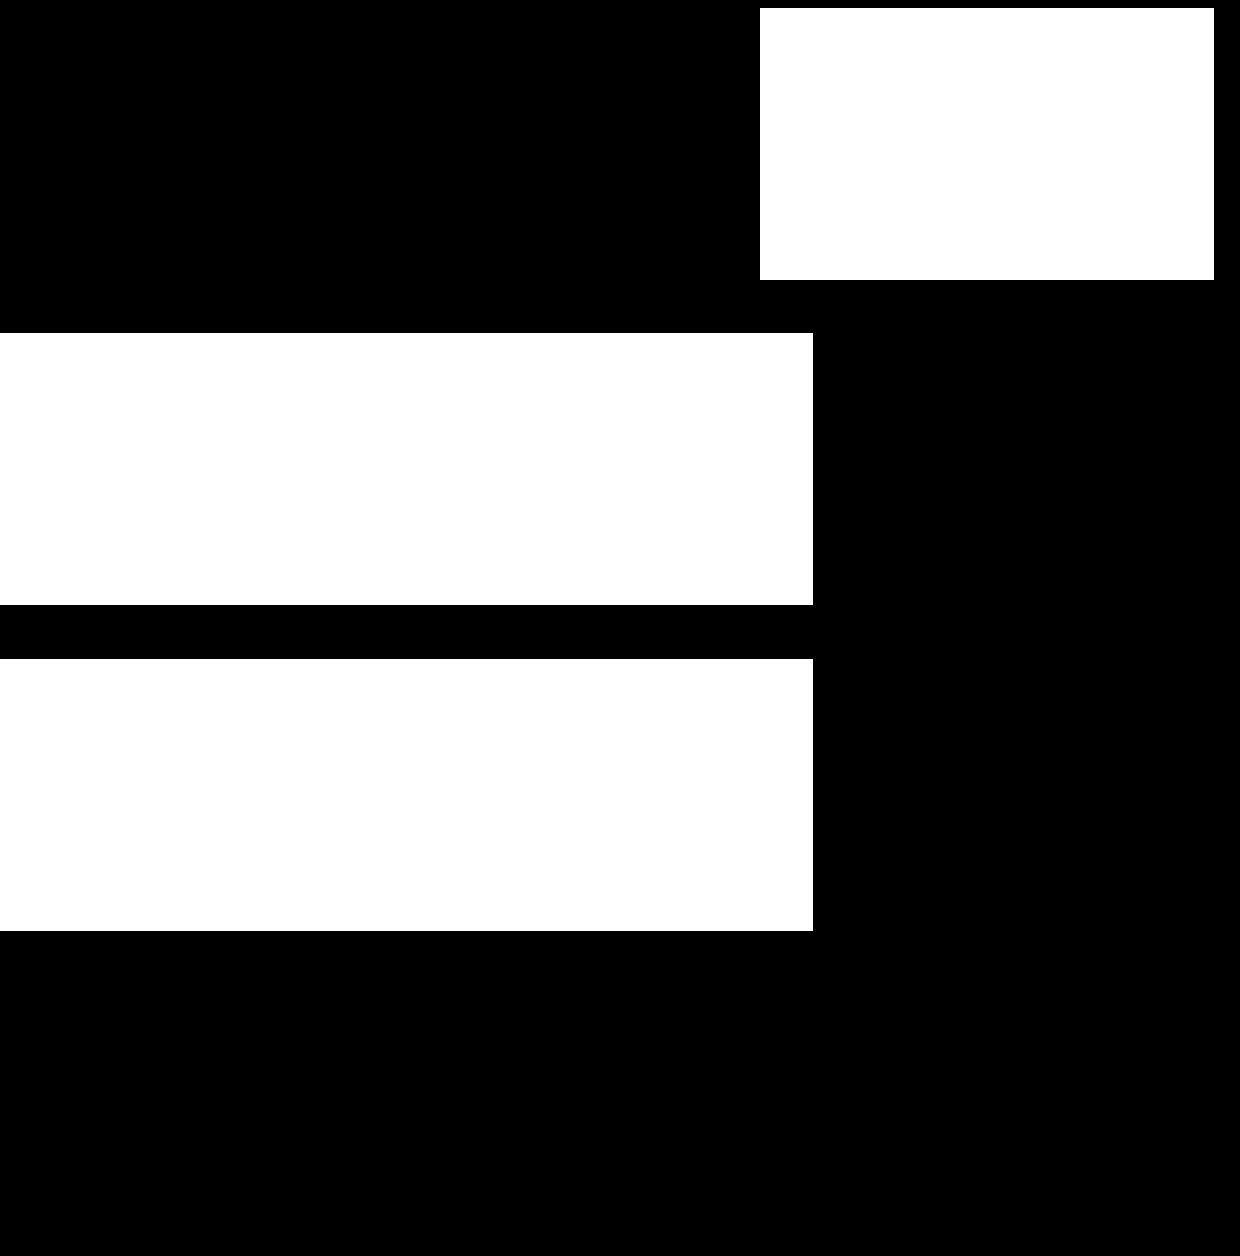
# The codebase: Meet the CLI



# The codebase: Load gen



# The codebase: The proxy



# Moving to a new data store

EXERCISE: [01]



# Game plan

EXERCISE: [01]

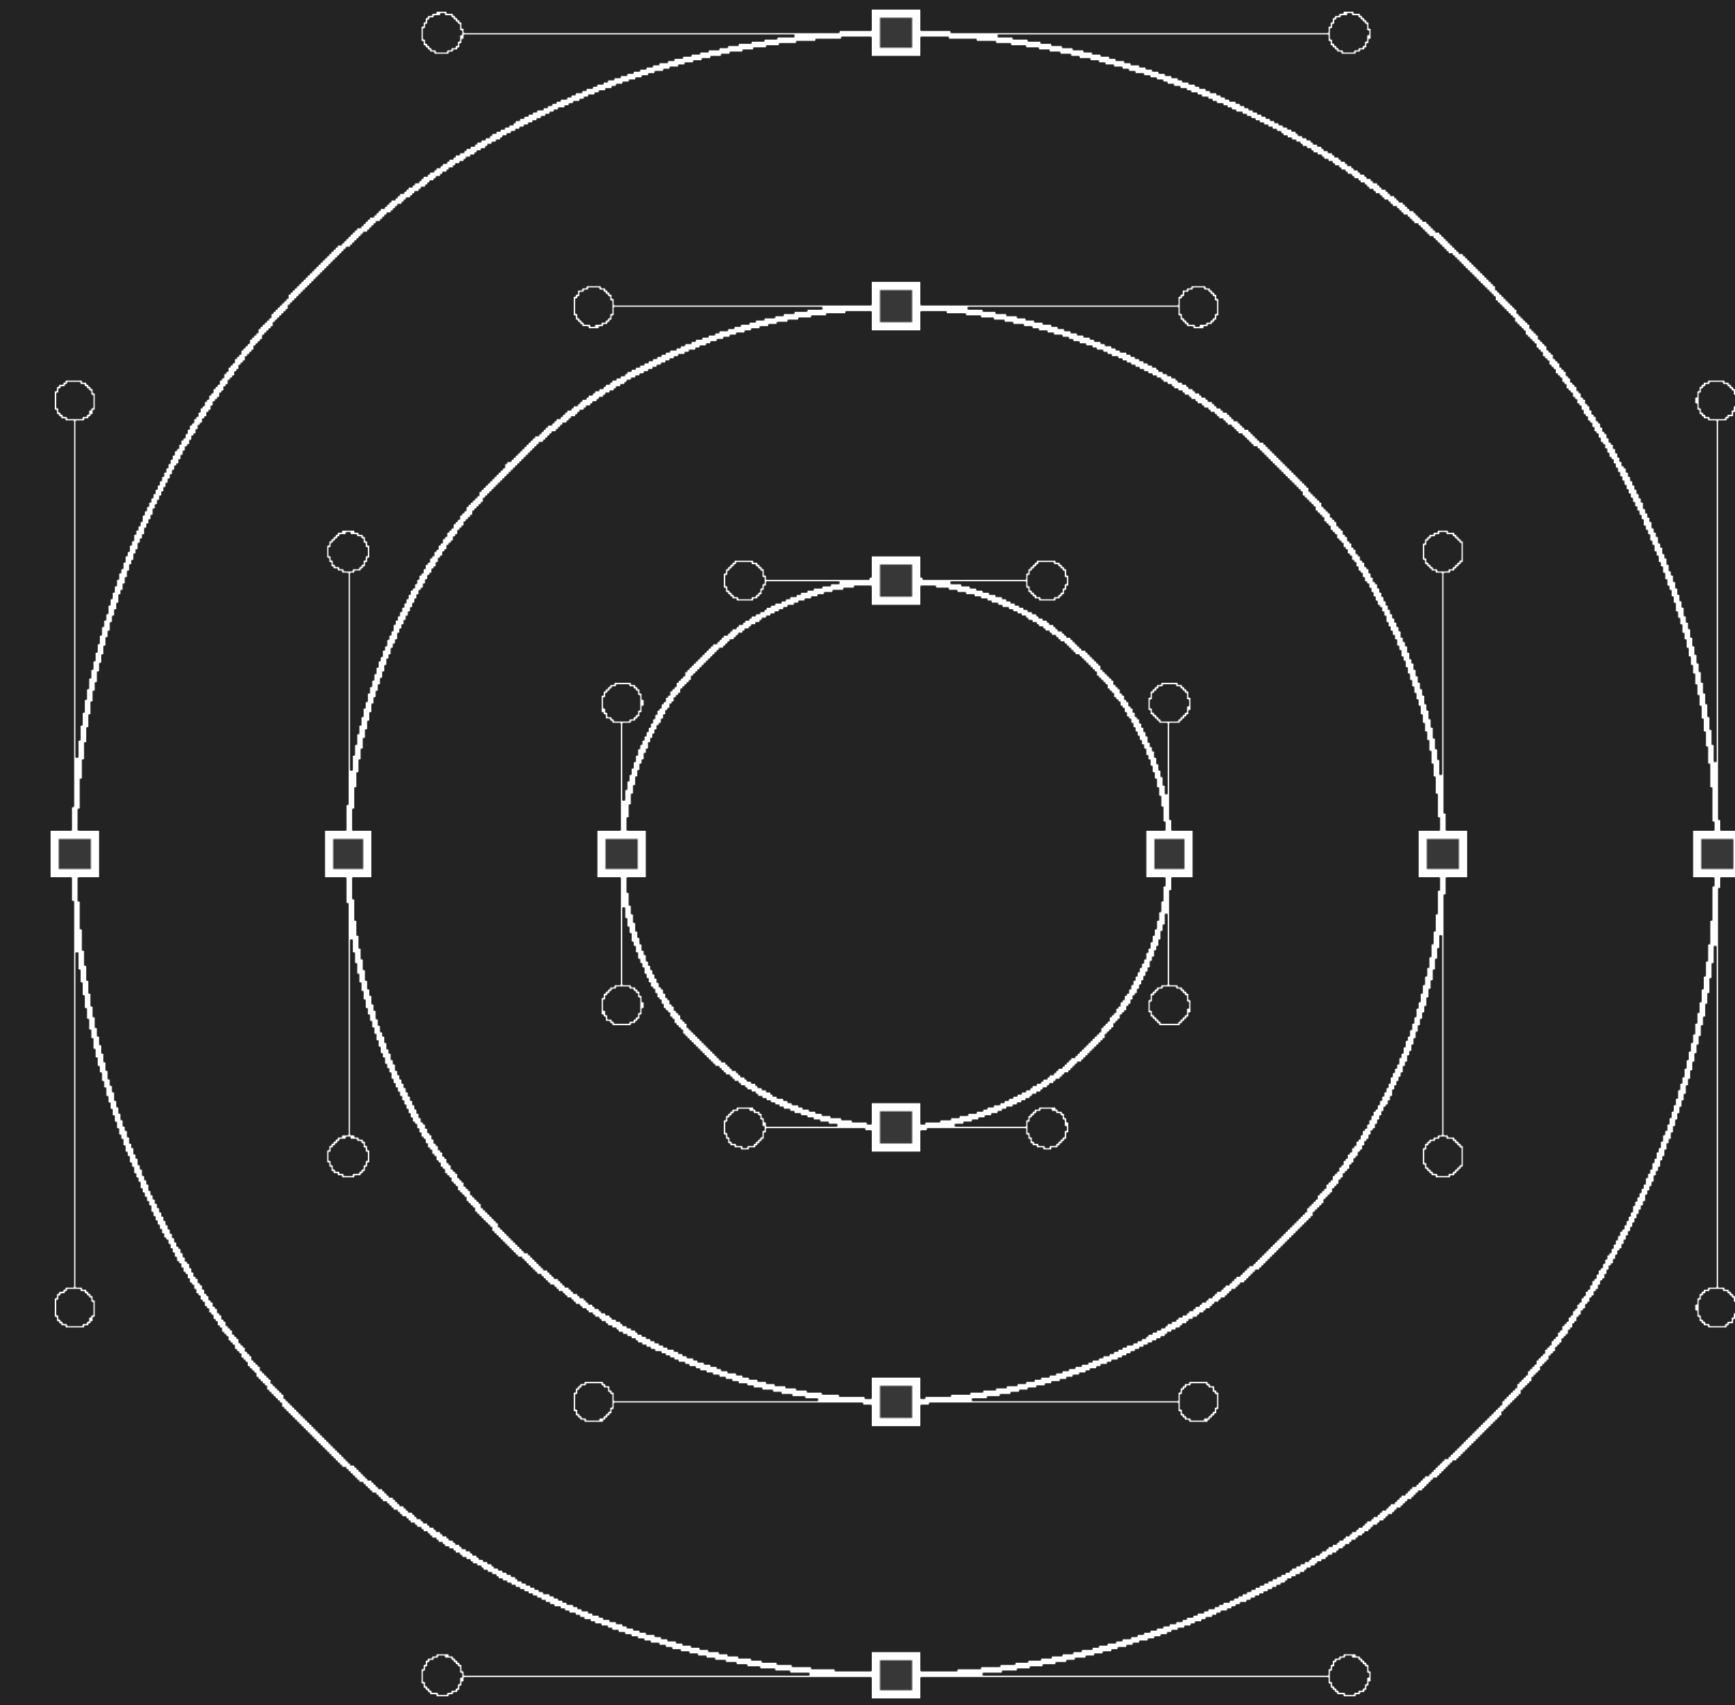


# Adding sharding

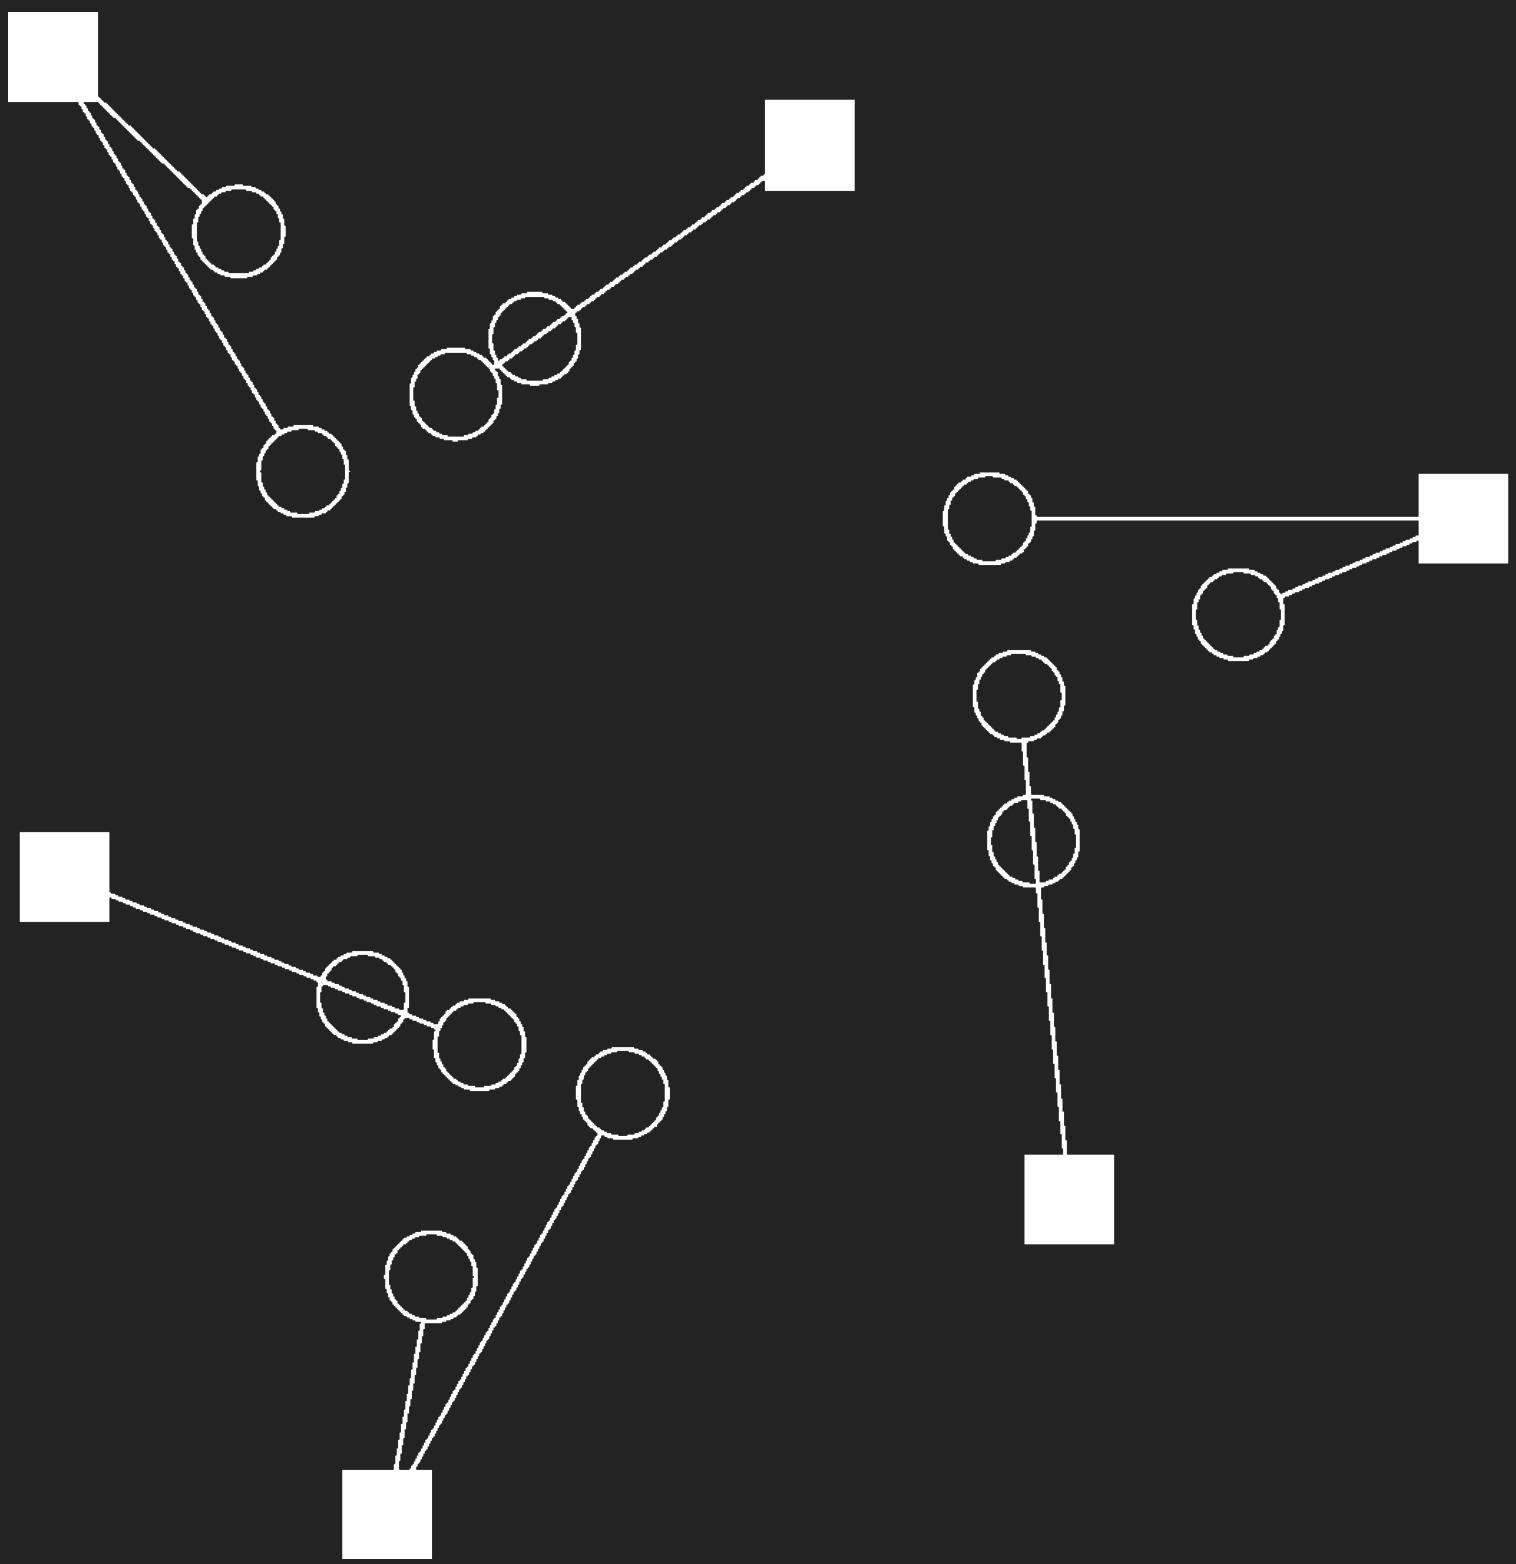
EXERCISE: [02]



# Wrapping up...



# What we learned today



## ■ What we learned today

- 
- 
- 

[01]

# Migrations are everywhere + necessary

## ■ What we learned today

- 
- 
- 

[02]

Migrations can be stressful,  
but with the right preparation,  
they don't have to be

## ■ What we learned today

- 
- 
- 

[03]

Can borrow from systems: atomic operations, observability, feature flags, traffic splitting, etc.

## ■ What we learned today

- 
- 
- 

[04]

Can plan migrations like systems:  
documentation/runbooks

## ■ What we learned today

- 
- 
- 

[05]

To ship faster, write better  
migrations

■ Thank you :)

- 
- ↳ Slides + code available
  - ↳ [bruno.sh/go25](http://bruno.sh/go25)
- 

BRUNO SCHEUFLER  
[AT]GOPHERCON  
UK → 2025



INNGEST  
E COM  
EST... 2021  
©2025



.

.

.

# Questions?



■ It takes a village <3

---

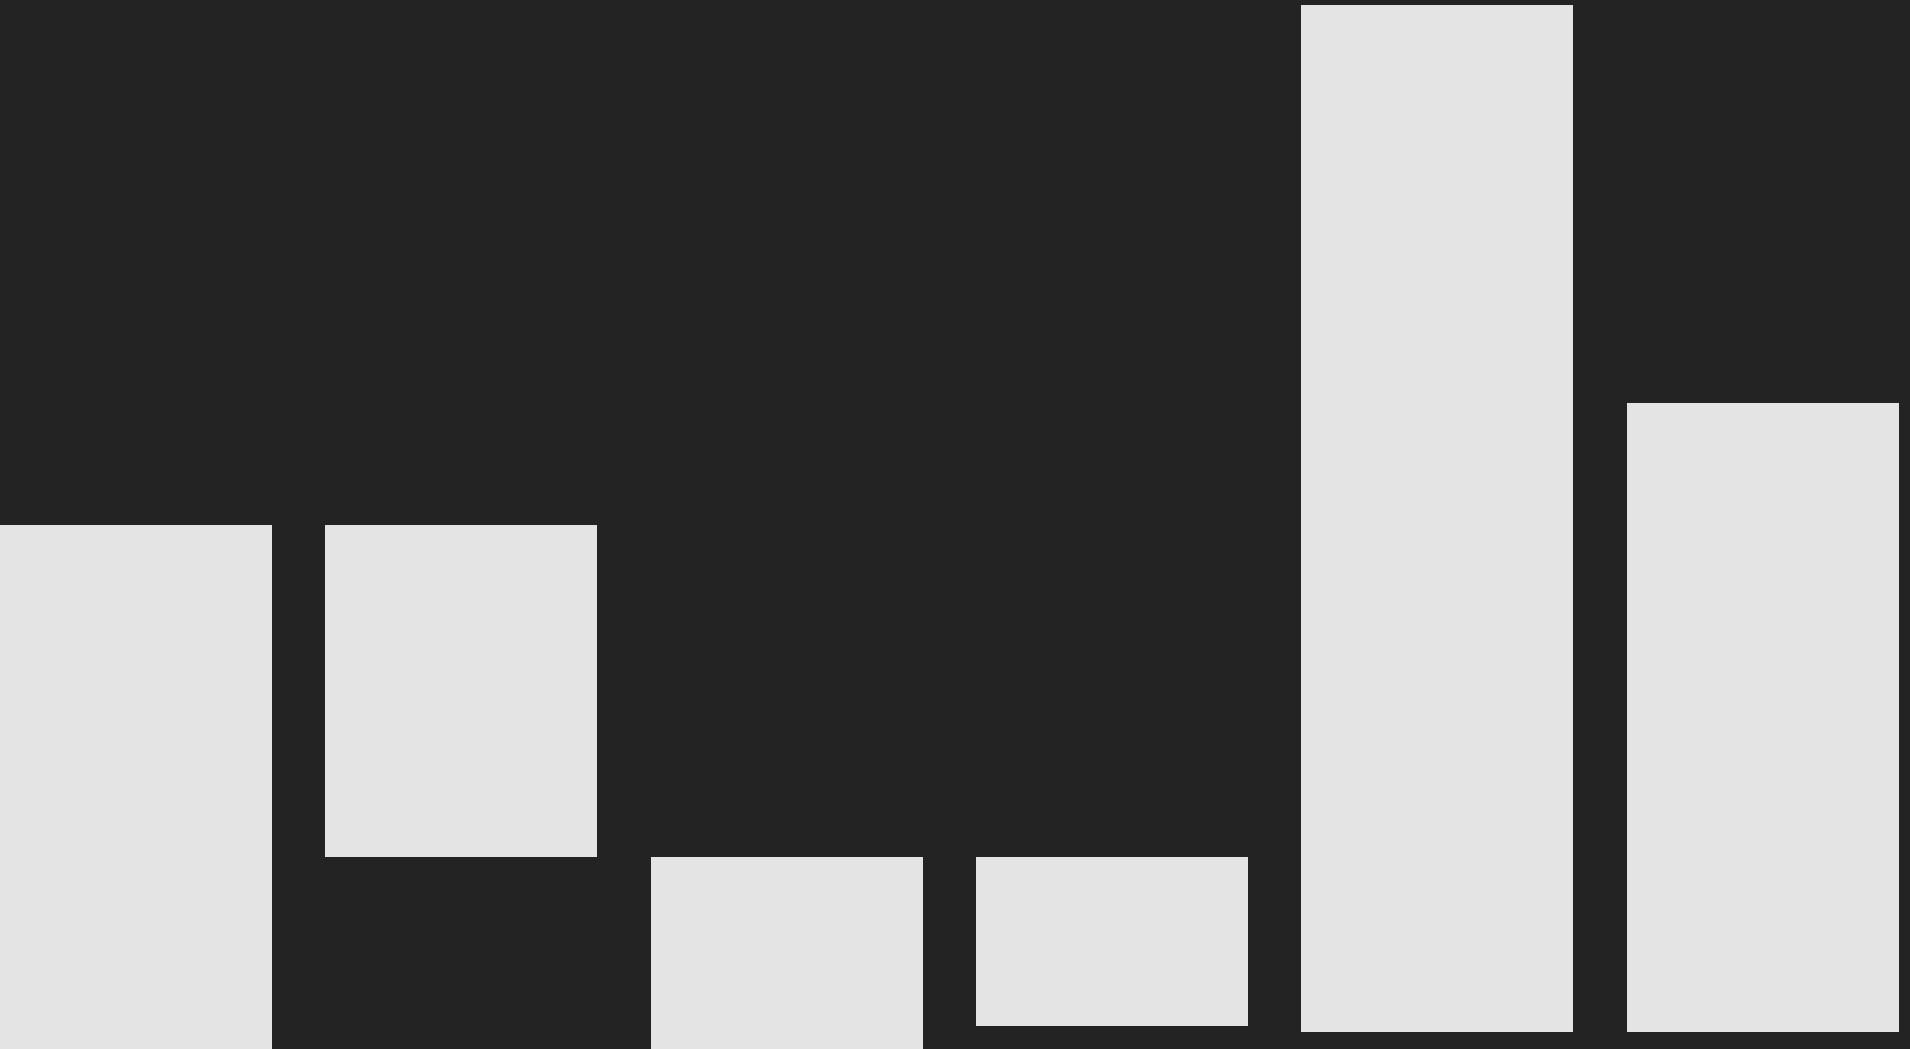
Thanks to all GopherCon UK organizers!

---

Design: Sophia Umansky  
↳ Brand Design / Inngest

---

Thanks to the amazing team at Inngest  
for doing everything to make this  
talk come to life.



BRUNO SCHEUFLER  
[AT] GOPHERCON  
UK → 2025  
SOFTWARE ENGINEER  
@INNGEST

INNGEST

INNGEST  
COM  
EST... 2021  
©2025

